



US006208180B1

(12) **United States Patent**
Fisch et al.

(10) **Patent No.:** **US 6,208,180 B1**
(45) **Date of Patent:** **Mar. 27, 2001**

(54) **CORE CLOCK CORRECTION IN A 2/N
MODE CLOCKING SCHEME**

4,293,927 10/1981 Hoshii 364/900
4,300,019 11/1981 Toyomaki 179/1 GE
4,365,290 12/1982 Nelms et al. 364/200

(75) Inventors: **Matthew A. Fisch; Chakrapani
Pathikonda**, both of Beaverton; **Javed
S. Barkatullah**, Portland, all of OR
(US)

(List continued on next page.)

FOREIGN PATENT DOCUMENTS

0103755 3/1984 (EP) G06F/1/00
0140814 5/1985 (EP) G06F/1/00

(List continued on next page.)

OTHER PUBLICATIONS

European Search Report, EP94306556, Aug. 24, 1995.
B. Case, "R400 extends R3000 architecture with 64-bit capabilities", Microprocessor Report vol. 5 Issue 19, pp. 10-14, Oct. 16, 1991.
Popescu, et al., "The Metaflow Architecture", IEEE Miro, pp. 10-13 & 63-73, 1991.
Kurpanek, et al., "PA7200: A PA-RISC Processor with Integrated High Performance MP Bus Interface", pp. 375-382, IEEE 1994.
Wilson, "MIPS rethinks RISC with superpipelining", Computer Design vol. 30, Issue 3, pp. 28-31, Feb. 1 1991.
Slater, "MIPS previews 64-bit R4000 architecture", Microprocessor Report, vol. 5, No. 2, pp. 1-7, Feb. 6, 1991.
Intel, "Automotive Components Handbook", 1988.
Burke, et al. "IBM shows tabletPC, compact desktop", Comdex-PC Week, Nov. 23, 1992.
Motorola, Inc. "R1", Motorola, INC., 1990.

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 0 days.

(21) Appl. No.: **09/170,997**

(22) Filed: **Oct. 13, 1998**

Related U.S. Application Data

(60) Division of application No. 08/709,379, filed on Sep. 6, 1996, now Pat. No. 5,834,956, which is a continuation-in-part of application No. 08/581,400, filed on Dec. 29, 1995, now abandoned.

(51) **Int. Cl.⁷** **H03L 7/00**

(52) **U.S. Cl.** **327/141; 327/142; 327/144;
327/147**

(58) **Field of Search** **327/141, 142,
327/144, 147**

(56) **References Cited**

U.S. PATENT DOCUMENTS

3,096,338 7/1963 Dehnert 260/299
3,623,017 11/1971 Lowell et al. 340/172.5
3,715,729 2/1973 Mercy 340/172.5
3,895,311 7/1975 Basse et al. 331/1
3,919,695 11/1975 Gooding 340/172.5
3,931,585 1/1976 Barker et al. 331/1 A
3,936,762 2/1976 Cox, Jr. et al. 329/104
4,077,016 2/1978 Sanders et al. 331/4
4,095,267 6/1978 Morimoto 364/200
4,143,418 3/1979 Hodge et al. 364/200
4,203,153 5/1980 Boyd 364/200
4,264,863 4/1981 Kojima 328/39

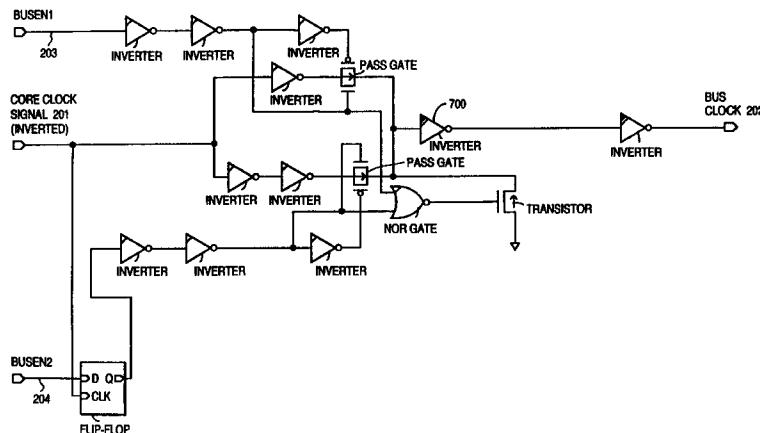
Primary Examiner—Margaret R. Wambach

(74) *Attorney, Agent, or Firm*—Blakely, Sokoloff, Taylor & Zafman LLP

(57) **ABSTRACT**

A 2/N mode dock generator that generates bus clock signals through the use of bus clock enable signals selecting bus clock pulses that are in phase and out of phase with a core clock signal. The clock generator maintains synchronization between the bus clock signal and the core clock signal so that they are always in a predetermined phase relationship.

15 Claims, 19 Drawing Sheets



U.S. PATENT DOCUMENTS


4,405,898	9/1983	Flemming	328/72
4,419,756	12/1983	Cheng-Quispe et al.	375/7
4,438,490	3/1984	Wilder, Jr.	364/200
4,479,191	10/1984	Nojima et al.	364/707
4,493,971	1/1985	Nawa et al.	219/137.5
4,615,005	9/1986	Maejima et al.	364/200
4,639,864	1/1987	Katzman et al.	364/200
4,667,289	5/1987	Yoshida et al.	364/200
4,669,099	5/1987	Zinn	377/47
4,698,748	10/1987	Juzswik et al.	364/200
4,758,945	7/1988	Remedi	364/200
4,766,567	8/1988	Kato	364/900
4,779,093	10/1988	Watkins	340/825.57
4,780,836	10/1988	Miyazaki et al.	364/551.01
4,780,843	10/1988	Tietjen	364/900
4,814,591	3/1989	Nara et al.	235/380
4,823,292	4/1989	Hillion	364/707
4,841,440	6/1989	Yonezu et al.	364/200
4,881,165	11/1989	Sager et al.	364/200
4,881,205	11/1989	Aihara	365/222
4,881,364	11/1989	Grenebach	56/314
4,896,260	1/1990	Hyatt	364/200
4,907,183	3/1990	Tanaka	364/707
4,922,450	5/1990	Rose et al.	364/900
4,931,748	6/1990	McDermott et al.	331/1 A
4,935,863	6/1990	Calvas et al.	364/138
4,967,895	11/1990	Speas	194/200
4,979,097	12/1990	Triolo et al.	364/200
4,979,190	12/1990	Sager et al.	375/106
4,980,836	12/1990	Carter et al.	364/483
4,991,129	2/1991	Swartz	364/707
5,003,537	3/1991	Sager	375/10
5,021,679	6/1991	Fairbanks et al.	307/66
5,025,387	6/1991	Frane	364/493
5,041,962	8/1991	Lunsford	364/200
5,059,924	10/1991	JenningsCheck	331/1 A
5,077,686	12/1991	Rubinstein	395/550
5,083,266	1/1992	Watanbe	395/275
5,103,114	4/1992	Fitch	307/271
5,117,443	5/1992	Shires	375/111
5,123,107	6/1992	Mensch, Jr.	395/800
5,125,088	6/1992	Culley	395/500
5,129,091	7/1992	Yorimoto et al.	395/750

5,133,064	7/1992	Hotta et al.	395/550
5,151,992	9/1992	Nagae	395/750
5,167,024	11/1992	Smith et al.	395/375
5,175,845	12/1992	Little	395/550
5,175,853	12/1992	Kardach et al.	395/650
5,191,657	3/1993	Ludwig et al.	395/325
5,220,672	6/1993	Nakao et al.	395/750
5,235,698	8/1993	Lan	395/550
5,239,631	8/1993	Boury et al.	395/325
5,239,652	8/1993	Seibert et al.	395/750
5,249,298	9/1993	Bolan et al.	395/70
5,251,320	10/1993	Kuzawinski et al.	395/750
5,256,994	10/1993	Langendorf	331/16
5,263,172	11/1993	Olnwich	395/800
5,276,824	1/1994	Skruhak et al.	395/375
5,287,384	2/1994	Avery et al.	375/1
5,305,452	4/1994	Khan et al.	395/550
5,319,771	6/1994	Takeda	395/550
5,319,772	6/1994	Hwang	395/550
5,336,939	8/1994	Eltrheim et al.	307/269
5,345,109	9/1994	Mehta	307/271
5,349,544	9/1994	Wright et al.	364/600
5,355,562	10/1994	Matoba et al.	24/625
5,359,232	10/1994	Eitrheim et al.	307/268
5,359,630	10/1994	Wade et al.	375/106
5,371,880	12/1994	Bhattacharya	395/550
5,381,542	1/1995	Carlson	395/550
5,388,250	2/1995	Lewis et al.	395/550
5,392,437	2/1995	Matter et al.	395/750
5,426,755	6/1995	Yokouchi et al.	395/425
5,428,765 *	6/1995	Moore	395/550


FOREIGN PATENT DOCUMENTS

0242010	10/1987	(EP)	G06F/1/00
0368144	11/1988	(EP)	G06F/1/32
0349762	1/1990	(EP)	B64D/11/06
0366250	5/1990	(EP)	H04M/11/06
0375794	7/1990	(EP)	G06F/15/00
0419908	4/1991	(EP)	G06F/1/32
0478132	4/1992	(EP)	G06F/13/42
0645717	3/1995	(EP)	G06F/13/42
2010551	6/1979	(GB)	G06F/1/00
2130765	6/1984	(GB)	G06F/1/00

* cited by examiner



FIG_1



FIG_2




FIG 3




FIG 4




FIG 5A




FIG 5B


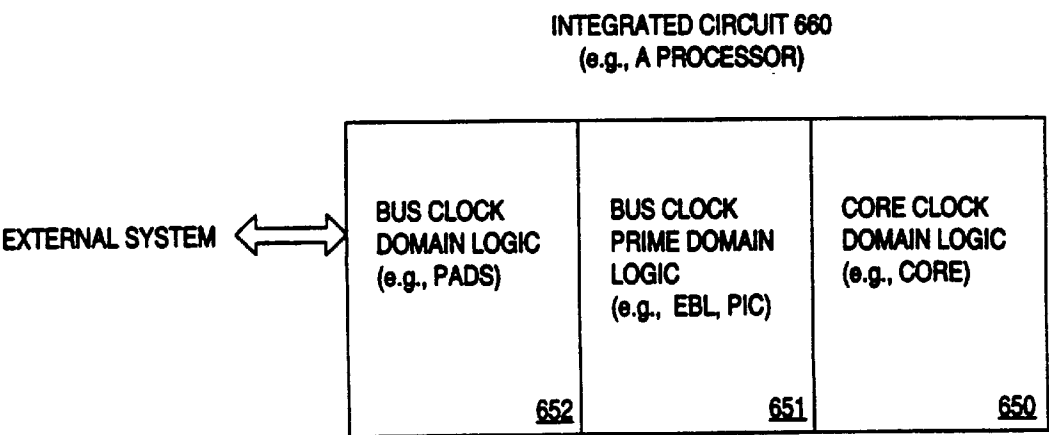


FIG 6A



FIG__6B


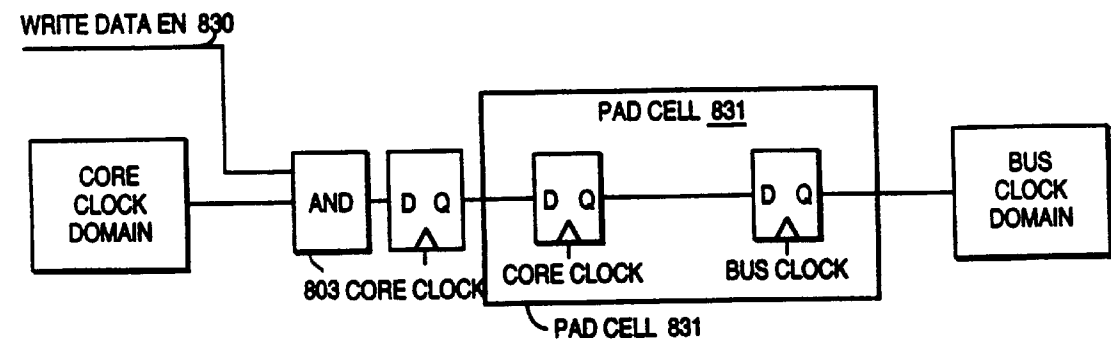

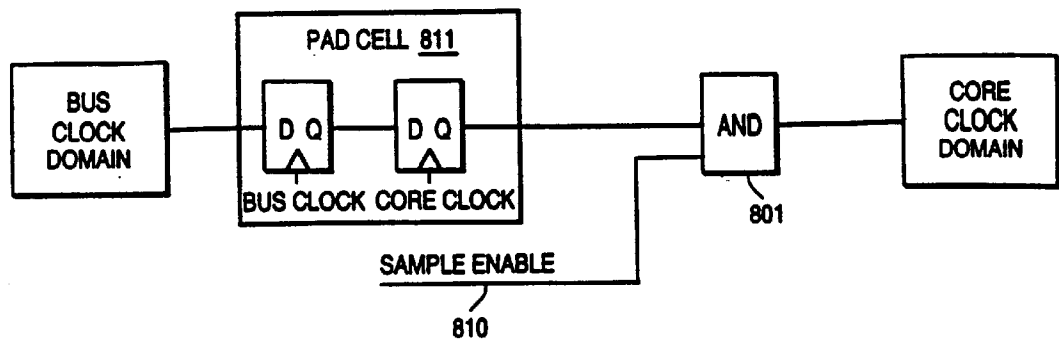



FIG. 7





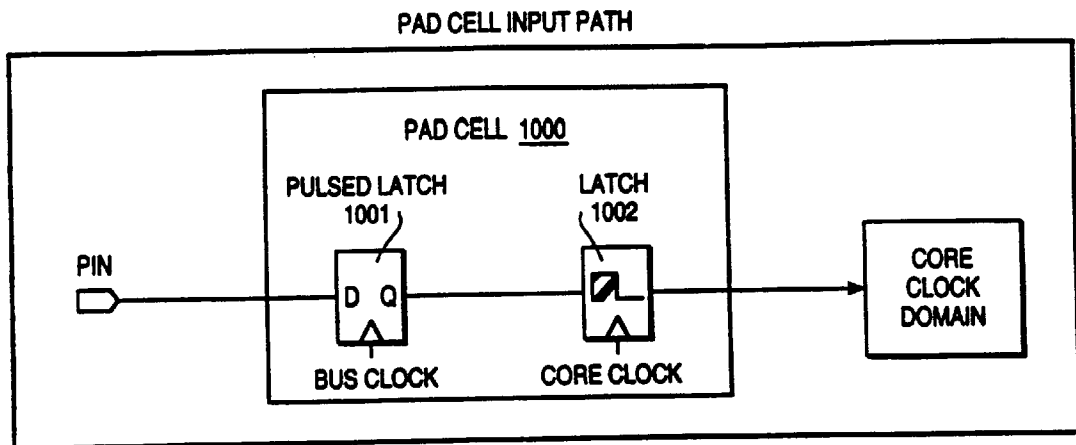


FIG. 10

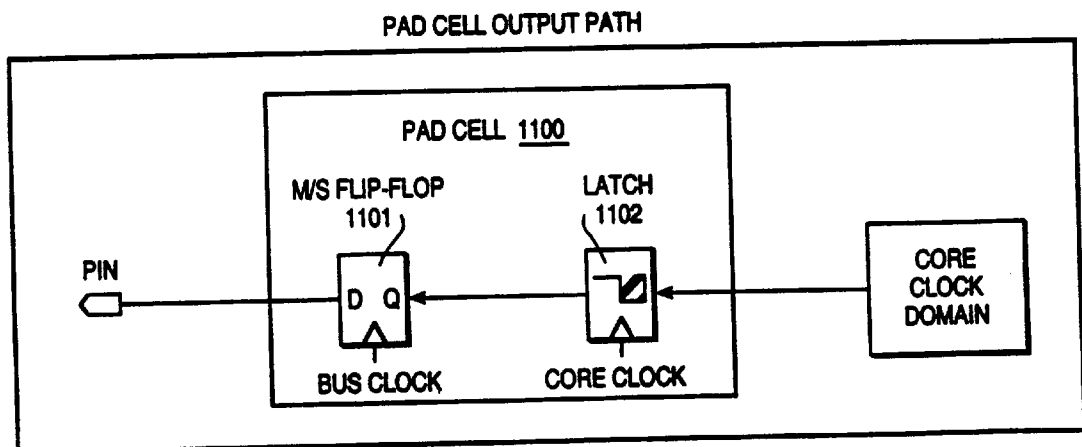
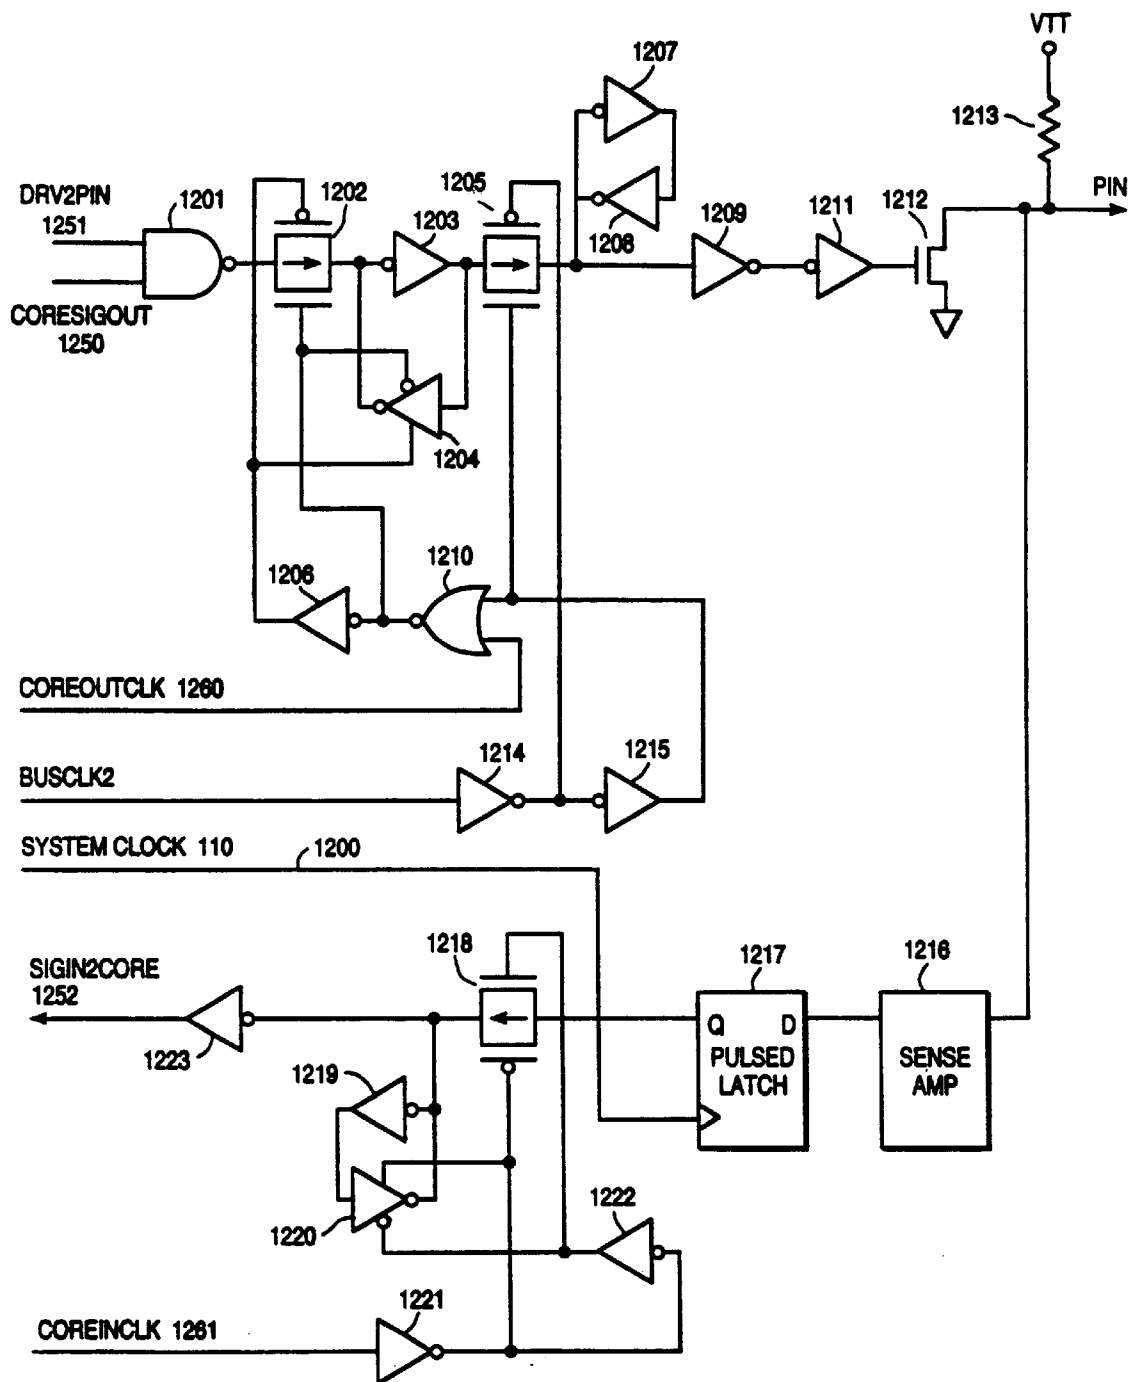


FIG. 11



FIG_12

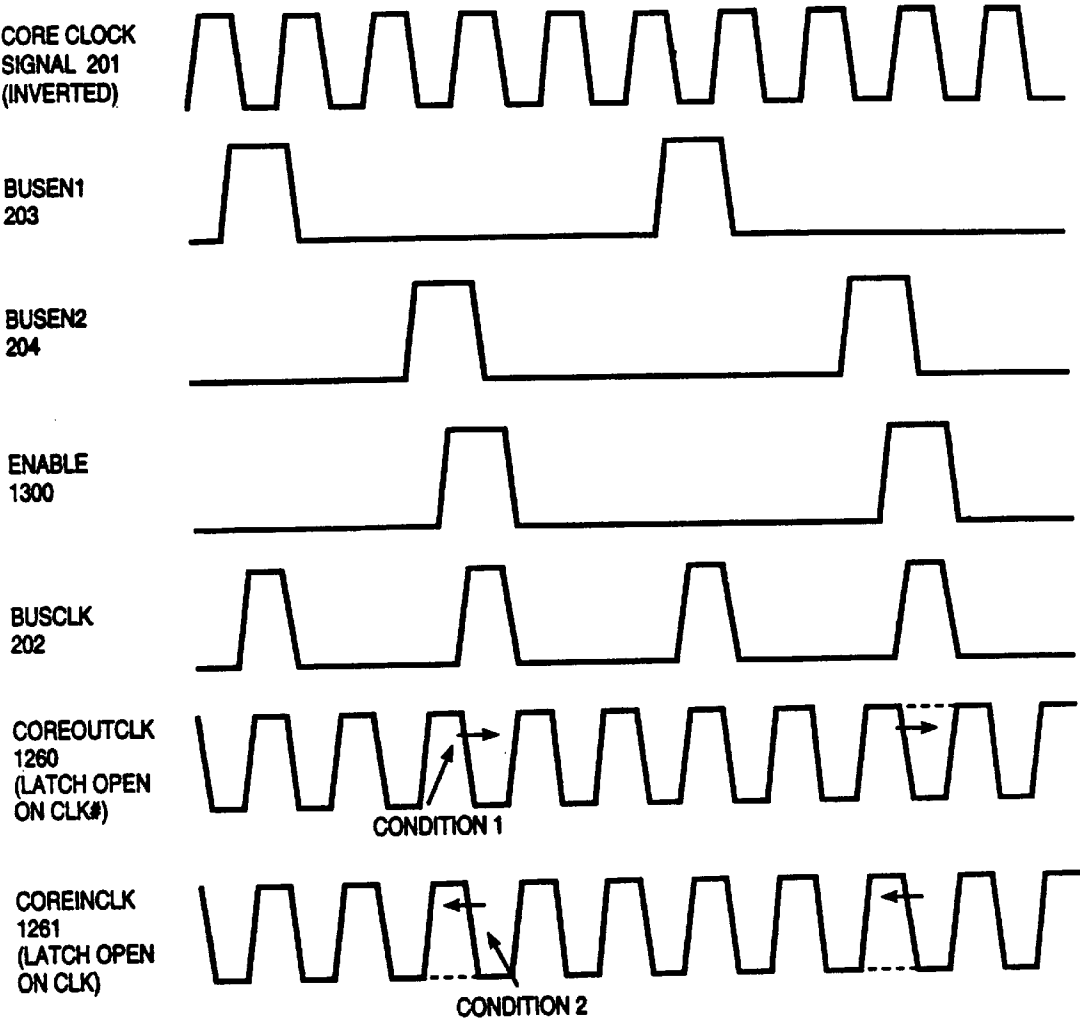


FIG 13

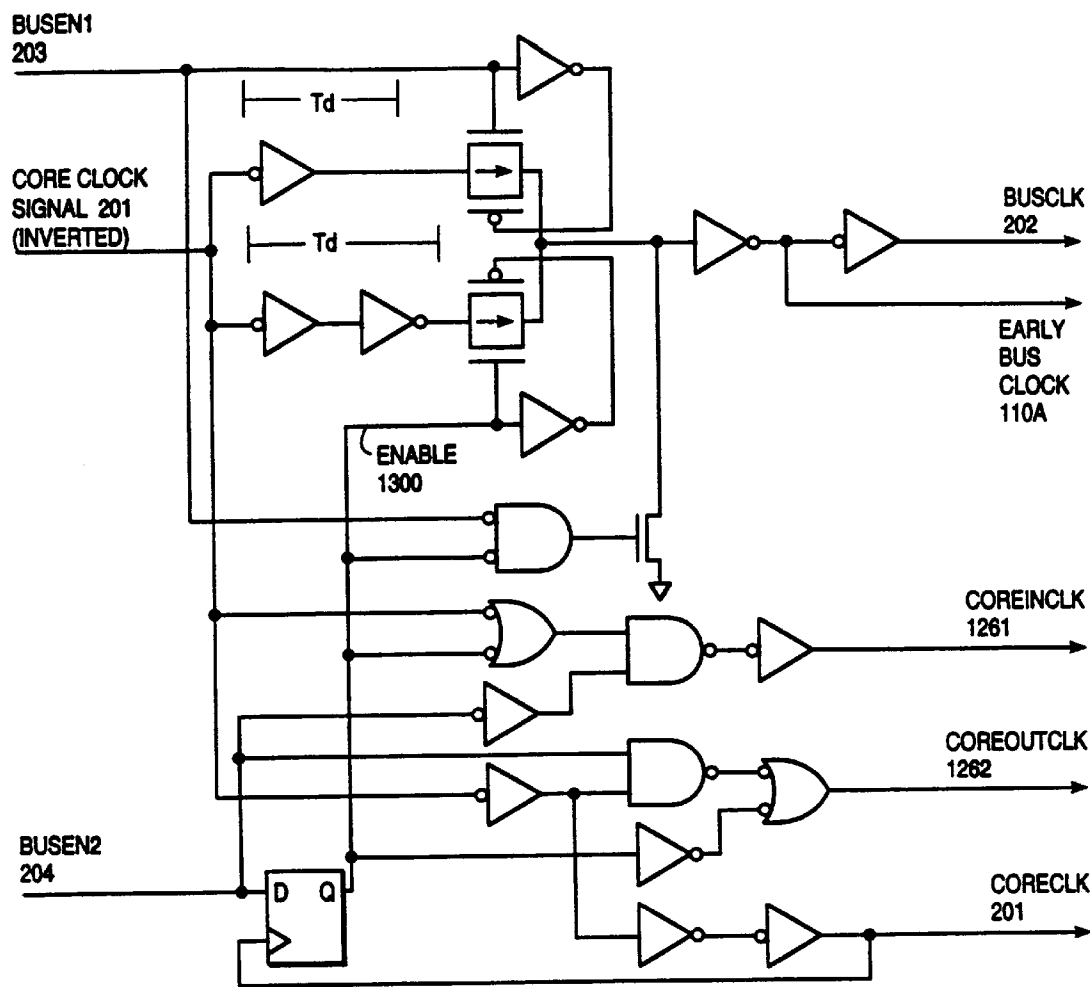


FIG. 14

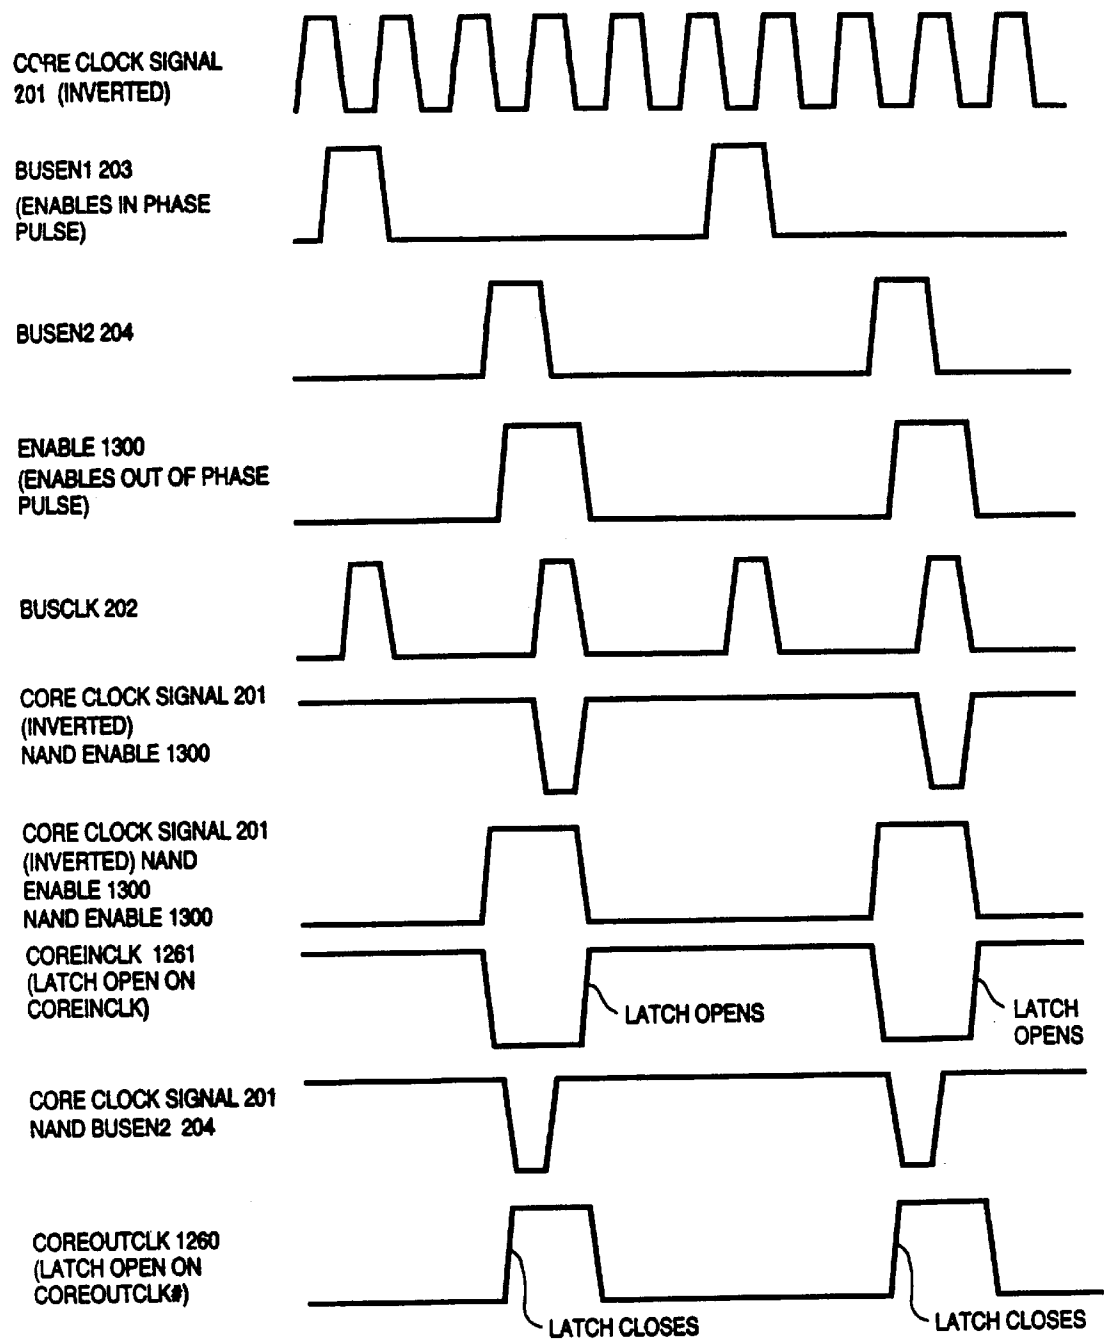


FIG 15

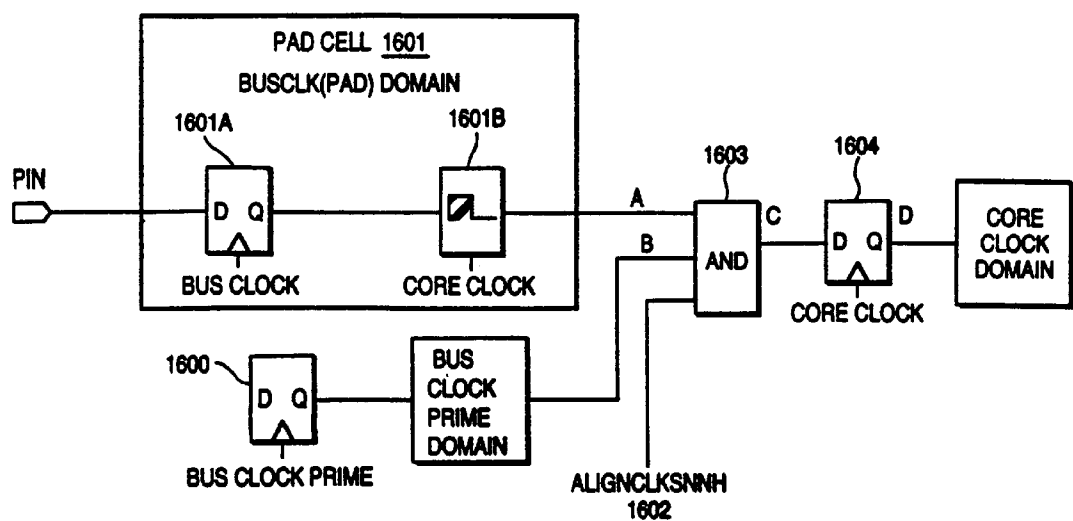


FIG 16

2/5 FRACTIONAL SPEED BUS




FIG 17

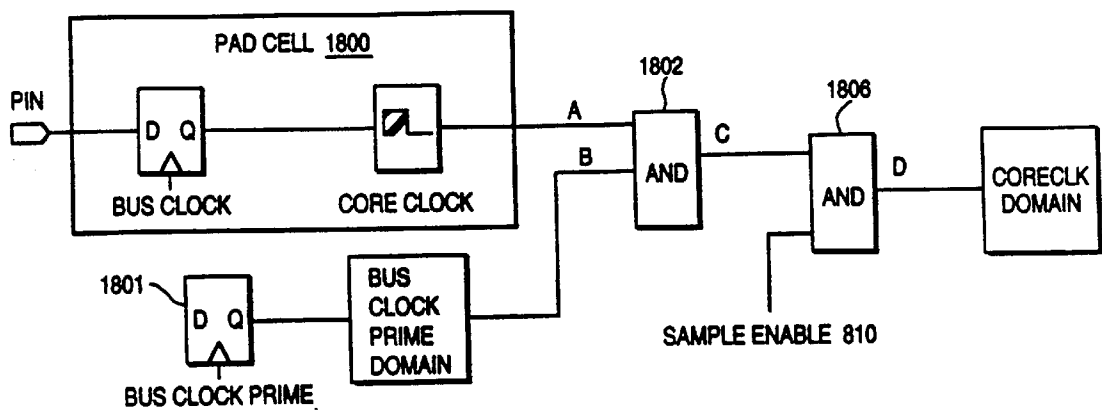


FIG __18A

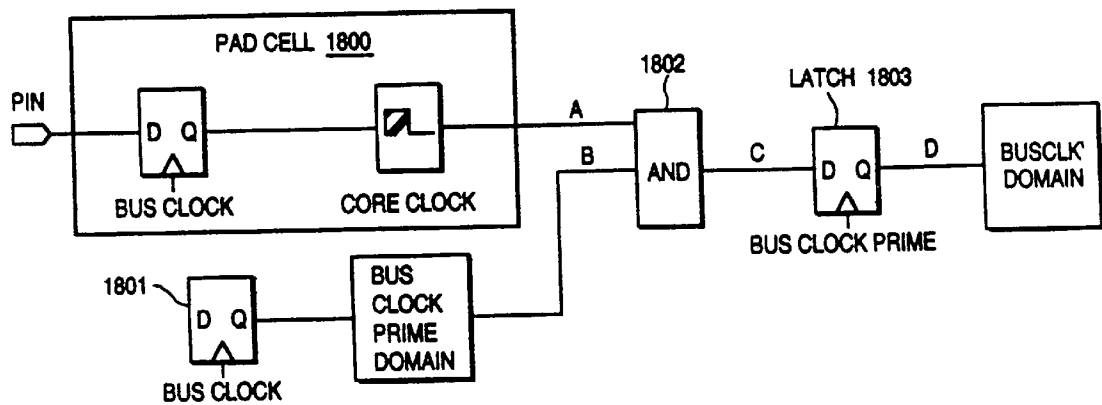





FIG __18B



FIG_19



FIG_20



FIG_21




FIG. 22A


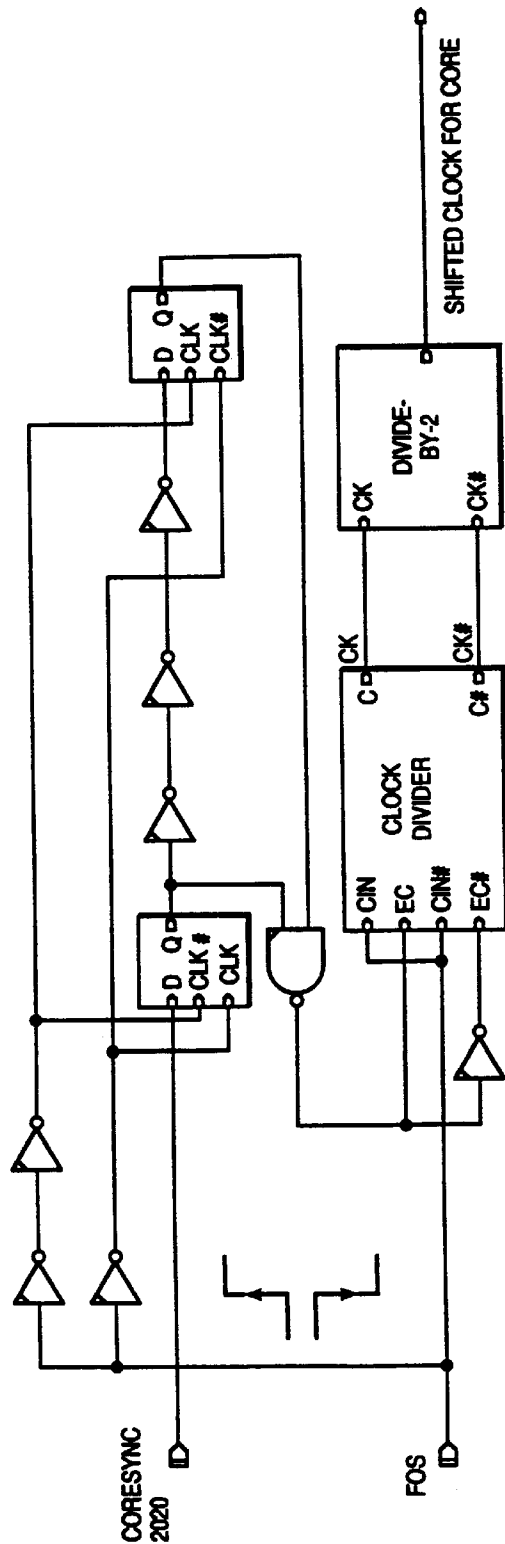
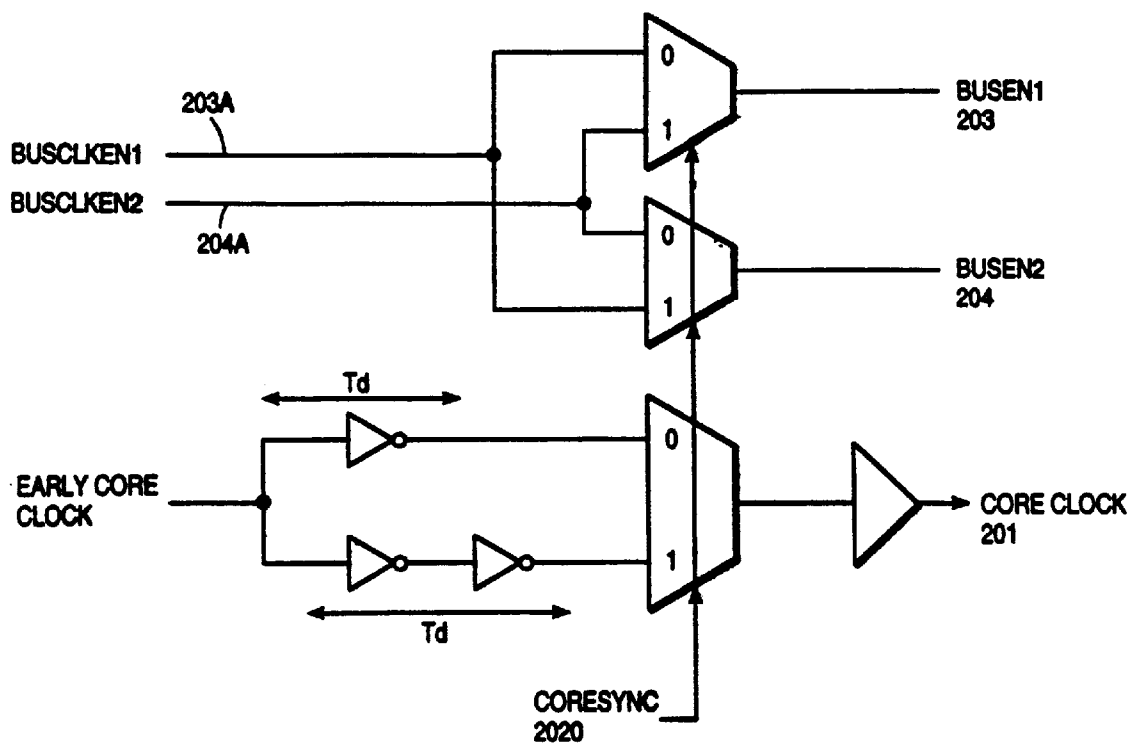
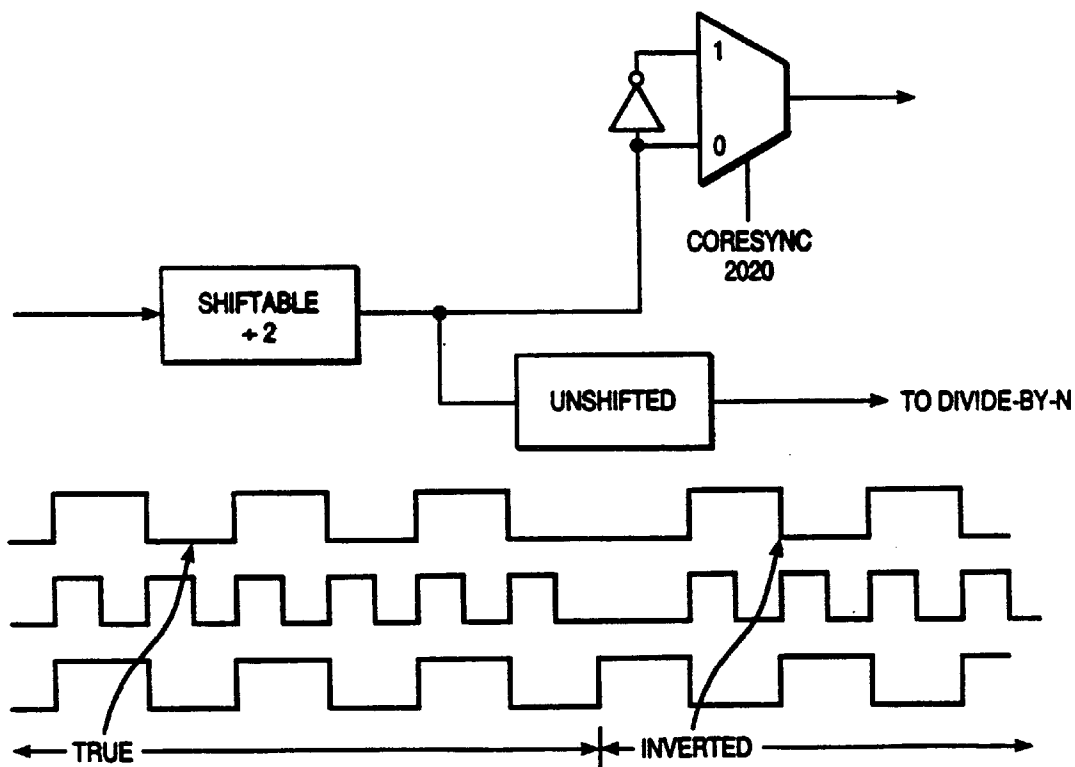


FIG. 22B



FIG_23



FIG_24

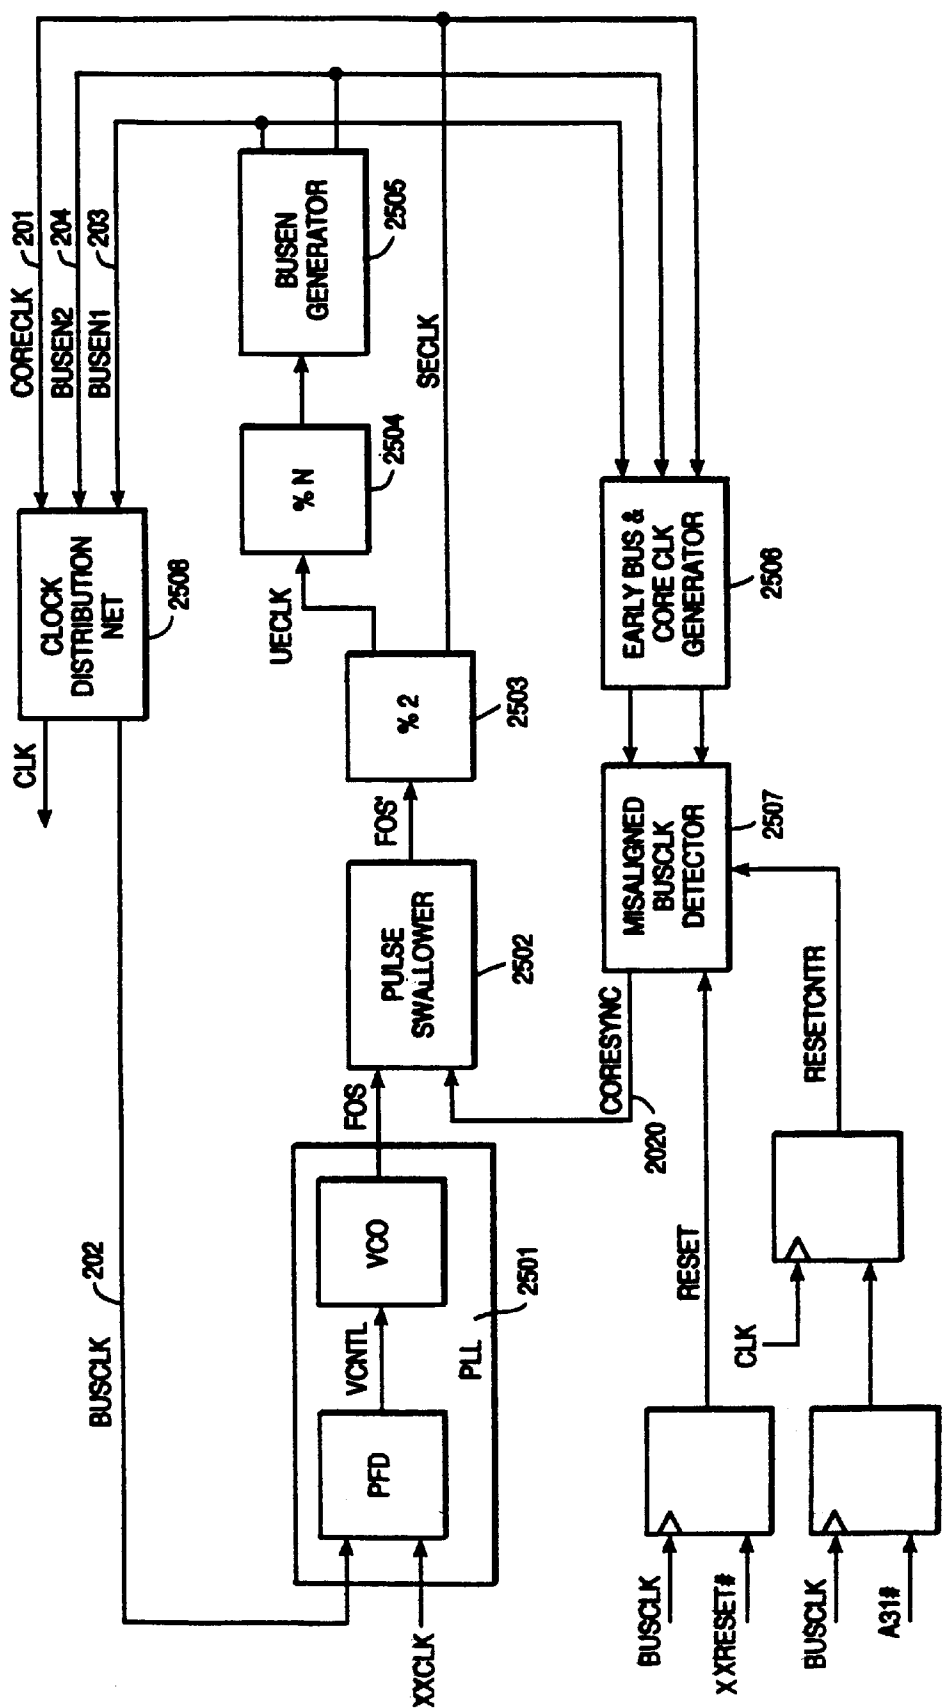


FIG. 25




FIG. 26

CORE CLOCK CORRECTION IN A 2/N MODE CLOCKING SCHEME

This is a continuation-in-part of U.S. patent application Ser. No. 08/581,400, filed Dec. 29, 1995, entitled "Method and Apparatus for Generating 2/N Mode Bus Clock Signals" now abandoned and a divisional of U.S. patent application Ser. No. 08/709,379, filed Sep. 6, 1996, now U.S. Pat. No. 5,834,956.

FIELD OF THE INVENTION

The present invention relates to the field of clock generation; more particularly, the present invention relates to generating bus clock signals and core clock signals, such as those having a 2/N ratio relationship with respect to each other.

BACKGROUND OF THE INVENTION

Historically the system frequency of a computer was limited by the available silicon technology. In other words, system board technology was capable of higher frequency operation than that of microprocessors provided by then available silicon technology. Computer systems designed using older technologies commonly operated both the system board and the microprocessor at the lower frequency required by the microprocessor. However, because silicon technology has advanced more rapidly over time than system board technology, a crossover has occurred. Today, advanced silicon technology provides microprocessors that are capable of operating at much higher frequencies than currently available system boards.

Besides system board and microprocessor technologies, other situations exist where differing technology constraints may result in distinct frequency domains. For example, a microprocessor and a cache memory may be manufactured using different silicon technologies, each technology providing a different maximum frequency of operation. A microprocessor and a floating point unit may also be operated advantageously at different frequencies. In these examples, the cache memory and/or floating point unit may be either on the same chip (or die) as the CPU, or on a separate chip. Even when manufactured on the same chip, individual functional blocks that exchange data may be advantageously operated at a different frequencies. Other examples of devices that may be advantageously operated at different frequencies are discrete logic components and telecommunications devices.

Today, in order to provide high performance computer systems, it is advantageous to operate the microprocessor(s) at its highest possible frequency and other parts of the computer system at a lower frequency dictated by either system board technology or other technology constraints. Transferring data between various components that operate at different frequencies may be synchronous or asynchronous.

Synchronous designs in microprocessors are advantageous over asynchronous designs for several reasons. First, microprocessor design validation tools are optimized for debugging synchronous logic designs. Using synchronous design techniques makes microprocessor design and validation much easier and more reliable. Second, synchronous designs are mandated in fault tolerant computer systems that use a master-checker scheme. In a master-checker system a first processor, the master, operates lock-stepped with a second processor, the checker. During each clock cycle, the checker processor monitors whether the two processors

produce identical results on their pins to provide error checking. Because asynchronous data transfer designs do not provide the requisite lock-stepped operation they are not acceptable in master-checker systems.

Some microprocessors operate internally at an integer multiple of the frequency of the computer system bus. In such a computer system, synchronous data transfer between components operating at different frequencies is facilitated by the fact that a data transfer edge of the lower frequency clock corresponds to a data transfer edge of the higher frequency clock. Therefore, the data transfer edge of the lower frequency clock (and/or corresponding data transfer edges of the higher frequency clock) can be used to cause data transfer. However, integer multiple frequency designs only take advantage of quantum-leap improvements in silicon technology. For example, if system board technology is limited to 50 MHz operation, then microprocessors operating at 100, 150, 200, . . . MHz can be used in an integer multiple design. Suppose that currently available silicon technology provides microprocessors capable of operating at 120 MHz. The additional 20% performance beyond 100 MHz is not utilized in an integer multiple design. It would be advantageous to provide a microprocessor design capable of operating at more flexible frequency ratios to the system bus while also providing synchronous data transfer between the microprocessor and other devices on the computer system bus. In general, it would be advantageous to provide synchronous data transfer between a variety of digital logic and memory devices that operate according to flexible operational frequency ratios. The present invention provides these advantageous results.

One prior art 1/N mode bus clock generation scheme supports 1/N ratio bus clock to core clock, where N is limited to being 2, 3, or 4. In 1/N mode clock scheme, a phase-locked loop (PLL) generates a core clock signal which is at N times higher frequency than the system clock received by the integrated circuit. A programmable ring counter generates an enable signal every Nth cycle which selects every Nth cycle of the core clock to generate the internal bus clock signal.

One advantage of such a scheme is that the core clock and bus clock share the same clock distribution network for many stages. This greatly reduces the possible skew between the core clock and the bus clock. In fact, the skew between the bus clock and the core clock is no worse than the intra-unit skew between core clocks. Thus, it is desirable to have a single clock distribution path for both the core clock and the bus clock, is particularly advantageous, allowing easy management of the skew between the two clocks.

At least one prior art microprocessor uses a phased locked loop (PLL) with some extra logic to generate a 2/3 mode clock. Although this uses the same PLL for bus clocks and core clocks, there is a separate clock distribution path for the two clocks.

Therefore, a system, method, and apparatus for synchronous data transmission between digital devices operating at frequencies having a 2/N integer ratio relationship is needed.

SUMMARY OF THE INVENTION

The present invention provides a clock generator fabricated on an integrated circuit is described. In one embodiment, the clock generator comprises a core clock generator configured to generate a core clock signal, a bus enable generation logic configured to generate first and second bus clock enable indications. A first circuitry is configured to generate a bus clock signal by selecting every

N/2 cycles of the core clock signal in response to the first and second bus clock enable indications. In one embodiment, N is not an integer greater than 2. The clock generator also includes a detector that is configured to determine whether the core clock signal is out of phase with the bus clock signal and correction circuitry that places the core clock signal in phase with a bus clock signal if the detector determines that the core clock signal is out of phase with the bus clock signal.

BRIEF DESCRIPTION OF THE DRAWINGS

The present invention will be understood more fully from the detailed description given below and from the accompanying drawings of various embodiments of the invention, which, however, should not be taken to limit the invention to the specific embodiments, but are for explanation and understanding only.

FIG. 1 is a block diagram of the embodiment of an integrated circuit of the present invention.

FIG. 2 illustrates a set of 2/N mode clock enable signal waveforms.

FIG. 3 illustrates one embodiment of a circuit to generate a bus clock signal using two bus clock enable signals.

FIG. 4 illustrates one embodiment of the enable signal generation block of the present invention.

FIGS. 5A and 5B illustrate a scheme for separating two bus clock enable signals from a combined bus enable signal.

FIG. 6A is a waveform diagram illustrating the bus clock prime signal.

FIG. 6B illustrates one embodiment of an integrated circuit having a bus clock prime domain.

FIG. 7 is an exemplary circuit implementation for generating the bus clock signal.

FIG. 8A illustrates the use of the sample enable signal.

FIG. 8B illustrates the use of the drive enable signal.

FIG. 8C illustrates the use of the write data enable signal.

FIG. 9 illustrates the clock and synchronization signals for the 2/5 fractional speed bus.

FIG. 10 illustrates the pad cell input path between the core clock domain and the pins.

FIG. 11 illustrates the pad cell output path between the core clock domain and the pins.

FIG. 12 illustrates one embodiment of a pad cell implementation for the 2/N fractional bus.

FIG. 13 is a waveform diagram illustrating mindelay problems with the two core clock signals.

FIG. 14 illustrates one embodiment of a clock generation circuit for generating the core clock signals.

FIG. 15 illustrates output waveforms of the distribution network of FIG. 14.

FIG. 16 illustrates one embodiment of a circuit for handling 2/N glitch.

FIG. 17 is a timing diagram illustrating a glitch for the 2/5 fractional speed bus.

FIG. 18A illustrates an alternate embodiment for handling 2/N glitch.

FIG. 18B illustrates another embodiment for handling 2/N glitch.

FIG. 19 illustrates a waveform of a shifted pulse that occurs during correction.

FIG. 20 illustrates one embodiment of a detector logic diagram for a sync signal generator of the present invention.

FIG. 21 illustrates one embodiment of associated waveforms for a sync signal generator.

FIG. 22A illustrates an example of the 2x clock having one of its pulses swallowed.

FIG. 22B illustrates one embodiment of the circuitry and associated waveforms for shifting the core clock phase.

FIG. 23 illustrates pulse shifting circuitry using an early core clock signal.

FIG. 24 illustrates one embodiment of core signal correction logic of the present invention.

FIG. 25 is a block diagram of one embodiment of a 2/N clocking circuit of the present invention.

FIG. 26 is a timing diagram of the bus clock enable signals before and after a shift of the core clock signal.

DETAILED DESCRIPTION OF THE PRESENT INVENTION

The method and apparatus for generating clock signals in an integrated circuit is described. In the following detailed description of the present invention, numerous specific details are set forth, such as specific signal names, device frequencies, bus frequencies, 2/N frequency ratios, etc., in order to provide a thorough understanding of the present invention. However, it will be apparent to one skilled in the art that the present invention may be practiced without these specific details. In other instances, well-known structures and devices are shown in block diagram form, rather than in detail, so as not to obscure the present invention.

It is understood that the present invention may be comprised of transistor circuits that are readily manufacturable using well-known CMOS (complementary metal-oxide semiconductor) technology, or other equivalent semiconductor manufacturing processes. In addition, the present invention may be implemented with other manufacturing processes for making digital devices.

While in the following discussion the present invention is presented with respect to implementation in a microprocessor, the present invention is not limited to that implementation. Implementations for generating clock signals for various digital devices such as discrete logic devices, memory devices, devices either on the same or separate chips, communications devices, etc., are within the scope and spirit of the present invention.

The 2/N Mode Clocking Scheme

The present invention provides for a 2/N mode clocking scheme. In such a scheme, for every N core clocks, there will be two bus clock cycles. Again, N may be any integer value, for example, 5, 7, and 9, to create such ratios as 2/5, 2/7, or 2/9, respectively, or for example, 4, 6, 8, etc., to create ratios, such as 2/4, 2/6 and 2/8, respectively. Specifically, in one embodiment, the core operates in conjunction with one or more core clock signals having a frequency of 200 MHz, while transfers with the bus occur using bus clock signals having a frequency of 66 MHz.

In one embodiment, the present invention provides a scheme for generating bus clock signals and core clock signals for an integrated circuit, where the bus and core clock signals have a 2/N ratio relationship with respect to each other, with N being an odd integer 3 or greater (e.g., 5, 7, 9, etc.). In other words, the frequency of core clock signal is N/2 times higher than the frequency of the bus clock signal. Note that N may be an even integer; however, in such a case, the core clock signal is an integral multiple of the bus clock signal. The present invention provides an enhancement over 1/N methodology, even where N in the present

5

invention is an even integer value. In the present invention, the core clock signal refers to the clock that runs the internal portion of a device (e.g., processor), while the bus clock signal refers to the clock used for clocking bus operations, including transfers over a bus to and from the core.

FIG. 1 is a block diagram of one embodiment of an integrated circuit of the present invention. Referring to FIG. 1, the integrated circuit 100 is shown having a core 101, an interface 105, and a clock generator 102. The core 101 includes circuitry and logic to perform the designated functions of the integrated circuit, while the interface 105 provides an interface between the core 101 and the remainder of the system and its system bus(es). For instance, if the integrated circuit 100 comprises a processor, the core 101 may include one or more decoders, scheduling logic, execution units, reorder buffers, memory order buffers, register files, cache memory, etc., for use in executing instructions. The interface 105 may comprise external bus controller logic and programmable interrupt controller logic.

The clock generator 102 generates the clock signals in response to a system clock signal 110. In one embodiment, the clock generator 102 comprises a phase-locked loop (PLL). The clock signals are coupled to the core 101 and the interface 105. In one embodiment, the clock generator 102 generates the bus clock signal(s) 103 and the core clock signal(s) 104.

In one embodiment of the 2/N mode clock scheme of the present invention, the clock generator 102 generates the bus clock signal from the core clock signal using two bus clock enable signals, or other similarly functioning indications. These two bus clock enable signals are referred to herein as BusEn1 and BusEn2.

FIG. 2 illustrates a set of 2/N mode clock waveforms. Referring to FIG. 2, a core clock signal 201 and a bus clock signal 202 are shown. The bus clock signal 202 is generated by selecting cycles of the core clock signal 201 using the BusEn1 signal 203 and the BusEn2 signal 204. Thus, every Nth/2 cycle of the core clock signal 201 is selected to generate the bus clock signal 202. The BusEn1 signal 203 is used to select bus clock pulses which are in phase (i.e. aligned) with the core clock signal 201, while the BusEn2 204 selects bus clock pulses which are out of phase (i.e. misaligned) with the core clock signal 201.

As discussed below, in one embodiment, the present invention generates the bus clock enable signals using a programmable counter (e.g., a ring counter). Because both bus clock enable signals, the BusEn1 signal 203 and the BusEn2 signal 204, are generated by the same counter, they are always in synchronization with each other.

Note that in the case of N being an even integer, the 2/N mode is the same as a 1/N mode and only requires one of the two bus enables. In one embodiment, when 1/N mode is selected, the BusEn2 signal 204 remains low or is otherwise disabled.

FIG. 3 illustrates one embodiment of a circuit to generate the bus clock signal 202 using the BusEn1 signal 203 and the BusEn2 signal 204 in conjunction with an inverted version of the core clock signal 201. Referring to FIG. 3, the circuit comprises inverters 301–303 and NAND gates 304–306. Inverter 301 is coupled to receive the core clock signal 201 and inverts the core clock signal 201. The inverted core clock signal is coupled to one input of NAND gate 304. The other input of NAND gate 304 is coupled to the BusEn1 signal 203. The core clock signal 201 is also coupled to the input of inverter 302, which is coupled in series with inverter 303. The output of inverter 303 is coupled to one input of NAND gate 305, while the other input of NAND gate 305 is coupled to the BusEn2 signal 204.

6

Inverter 301 represents one inversion path, while inverters 302 and 303 together represent another. In the present invention, both inversion paths have the same delay. In other words, the delay on the core clock signal 201 caused by inverter 301 is the same as the delay caused by inverters 302 and 303 together (i.e., the delays are matched). The delays are the same to prevent jitter of the even and odd bus clock cycles relative to the system clock. Preventing jitter is important because the set up and hold time at the input/output (I/O) pins has to be increased by any delay difference and the jitter may cause PLL instability since the PLL tries to line up the edges of the bus clock signal to the edges of the system clock signal.

The outputs of NAND gates 304 and 305 are coupled to the inputs of NAND gate 306. The output of NAND gate 306 comprises the bus clock signal 202. In this manner, NAND gates 304–306 operate as output logic for the circuitry to generate the bus clock signal. It would be apparent to one skilled in the art that alternative logic gates or circuitry which performs the same functionality in response to the enable signals may be used.

Generation of the Enable Signals

In the present invention, the bus enable signals are generated by bus clock enable signal generation block in the processor. In one embodiment, the bus clock enable signal generation block supports ratios of 2/5, 2/7 and 2/9, as well as other ratios where N is even. In an alternative embodiment, other ratios such as 2/3 or greater than 2/9 may be supported.

FIG. 4 illustrates one embodiment of the enable signal generation block of the present invention. Referring to FIG. 4, a 3-bit latch 401 is coupled to receive the N value. In one embodiment, the N value is received from a pin. In an alternate embodiment; the N value may be received from a register (hardwired or programmed), a bond option, etc. Latch 401 is clocked by an active low (#) reset signal 406, such that whenever reset signal 406 is asserted, the value of N is latched into latch 401. In one embodiment, the reset signal 406 is asserted by another device in the system.

Latch 401 is coupled to a 3-to-8 decoder 403 and a fractional mode enable logic block 402. The fractional mode enable logic 402 latches in the value from 3-bit latch 401 in response to clock signal 407. The fractional mode enable logic block 402 determines if 2/5, 2/7 or 2/9 mode is selected and, if so, appropriately asserts a signal 410 to enable the BusEn2 signal 204. If the fractional mode enable logic block 402 determines that a 1/N mode has been selected, then the signal 410 is not asserted, thereby causing the BusEn2 signal 204 to be held low.

The 3-to-8 decoder 403 decodes the latched-in N value and supplies the decoded signals to a 9-bit ring counter 404. In another embodiment, an incrementer, state machine counter, shift register, etc., may be used to provide similar functionality to control the bus enable generational logic. Ring counter 404 strobes the bus enable output logic 405, causing the bus enable output logic 405 to output the BusEn1 signal 203 and the BusEn2 signal 204 with the appropriate timing.

In one embodiment, the bus enable output logic 405 may be implemented in a manner similar to FIG. 5B described below. Also, if a 1/N mode is selected, such bus enable output logic 405 may be simplified to only output a single enable signal.

In one embodiment, decoder 403 and ring counter 404 support the but to core clock ratios shown in Table 1:

TABLE 1

Bus to Core Clock Ratio Selection	
N Value	Ratio
000	1/2
001	1/3
010	1/4
011	1/2
100	1/2
101	2/7
110	2/9
111	2/5

In one embodiment, a single composite bus clock enable signal representing logical OR of the two bus enable signals **203** and **204** may be used. This composite bus clock enable signal may be ANDed with the core clock signal to produce a bus clock signal. Use of the composite bus clock enable signal reduces the routing overhead on the processor; however, it does not eliminate the space required for decoding the composite bus clock enable signal.

FIGS. **5A** and **5B** illustrate a scheme for separating two bus clock enable signals from a composite bus clock enable signal. Referring to FIG. **5A**, waveforms of the system clock signal **110**, the composite bus clock enable signal **511** and the select signal **510** are shown. The select signal **510** indicates which bus clock enable signal, BusEn1 or BusEn2, is being pulsed.

Referring to FIG. **5B**, the bus enable generation circuitry comprises a flip-flop clocked by the bus clock enable signal **511**. The data input of flip-flop **501** is coupled to an inverted version of the core clock signal **201**. On each pulse of the composite bus clock enable signal **511**, a clock cycle is output as a select signal **510** to an input of NAND gate **502**. The other input to NAND gate **502** is coupled to the bus clock enable signal **511**. The select signal **510** is also coupled to one input of NOR gate **505**. The other input to NOR gate **505** is coupled to the output of inverter **504**, which is coupled to receive the composite bus clock enable signal **511**. The output of NAND gate **502** is coupled to an input of inverter **503**. The outputs of inverter **503** and NOR gate **505** are the BusEn1 signal **203** and, BusEn2 signal **204**, respectively. The operation of the bus enable generation circuitry of FIG. **5B** would be within the understanding of those skilled in the art.

In 2/N mode, the falling edge of the composite bus clock enable signal **511** samples the high and low phase of the system clock signal **110** alternately. This sampling causes the select signal **510** to toggle every falling edge of composite bus clock enable signal **511**. The select signal **510** is then used to transfer alternate pulses of the bus clock enable signals BusEn1 and BusEn2 as outputs.

Thus, the present invention provides a simple scheme that does not need two phase lock loops (PLLs) to generate bus clock signals for a 2/N mode. In the prior art, two PLLs were required to generate a 2/N clocking scheme. However, it is difficult to keep two PLLs synchronized. In the present invention, because the two bus clock enable signals are generated from the same counter, they are synchronized. Therefore, the present invention is advantageous in that only one PLL is used and no complex mechanism or circuitry is required to maintain synchronization. Moreover, the present invention provides the bus clock and core clock signals using a single distribution path. This prevents a large skew from being created between the bus and core clock signals that is attributed to separate distribution paths as is in the prior art. In an alternative embodiment, separate distribution paths may be used to generate the core and bus clock signals.

Bus Clock Prime

In one embodiment, in order to simplify the external bus logic (EBL) for interfacing the core of the device to the system and the programmable interrupt controller (PIC) for controlling response by the device to received interrupts, the bus clock enable signals are slightly modified. In alternative embodiments, other interface units and circuitry of integrated circuits such as a processor may also be clocked with bus clock signals that are slightly modified. In this case, instead of clocking the EBL and the PIC with bus clock signal **202**, a modified version of bus clock signal **202**, referred to herein as the bus clock prime signal, is used.

The bus clock prime signal is always aligned with the core clock signal. The bus clock prime signal has the same high pulse as bus clock signal **202** when the bus and core clock signals are in phase, but every other high pulse of the bus clock prime signal is a core clock phase earlier than the bus clock pulses of bus clock signal **202** that are out of phase with the core clock signal **201**. In other words, the pulse of the bus clock signal **202** that corresponds to the low phase of the core clock signal **201** is advanced by a phase so that it now comes on the high phase of the core clock. In the 2/3 fractional bus mode, there is a difference of two cycles and three cycles between pulses, repeating in a 2, 3, 2, 3, etc. fashion, instead of a pulse of the bus clock signal **202** occurring every 2.5 core clock pulses. FIG. **6A** is a waveform diagram illustrating the bus clock prime signal **601**. Also shown in FIG. **6A** is a bus clock prime enable signal **602** which is used to select pulses of the core clock signal **201** for the bus clock prime signal **601**.

Using the bus clock prime signal of the present invention allows much of external bus logic which acts to interfere the core to the system bus to operate according to 1/N. In other words, using the bus clock prime signal, the present invention allows 1/N logic to be enabled to support 2/N mode. Because the 1/N logic may be used, existing 1/N clock distribution schemes may be used, with slight modification as described below.

FIG. **6B** illustrates an integrated circuit according to the teachings of the present invention. Integrated circuit **660** includes a core clock domain logic **650** which is clocked by the core clock signal. In one embodiment, the core clock domain includes the core of a processor or other integrated circuit device. The integrated circuit **660** also includes a bus clock domain logic **652** which is clocked by the bus clock signal. In one embodiment, the bus clock domain logic **652** includes pads in the pad ring that interface the integrated circuit **660** to the external system. Interfacing the bus clock domain logic **652** to the core clock domain logic is the bus clock prime domain logic **651**. In one embodiment, as described above, the bus clock domain logic may comprise the EBL and the PIC, as well as other elements. The bus clock prime domain logic **651** is implemented in 1/N logic, supporting the 2/N mode logic with 1/N logic.

The bus clock prime signal is used to facilitate the transfer of data to and from the core clock domain. When the core of the processor is interacting with the external system, it may be inputting or outputting data. When data is within the core, either before being sent or after being received, the core operates on the data in the core clock domain. Without using the bus clock prime signal, the present invention must provide circuitry to handle data transfers in both the low and high phases of the core clock signal. However, using the bus clock signal to control the pads of signals to and from the core, data transfers occur on the same phase (e.g., high phase) and the extra circuitry to accommodate transfers on both phases is not necessary. Note that in an alternate

embodiment of the present invention, data could be transferred on the low phase of the bus clock signal.

In one embodiment, the bus clock enable signals **203** and **204** are generated to select pulses of the system clock to create the bus clock prime signal **601**, and not the bus clock signal **202**. In such a case, the bus clock signal **202** may be obtained by delaying one of the bus clock enables for the bus clock prime signal **601**, thereby generating the appropriate enable for use in creating bus clock signal **202**. That is, in one embodiment, enables are generated to create the bus clock prime signal **601** and one of the enables is modified so that it correctly generates the bus clock signal **202**.

FIG. 7 is an exemplary circuit implementation for generating the bus clock signal **202** as mentioned above. Referring to FIG. 7, a circuit is shown receiving the BusEn1 signal **203**, the BusEn2 signal **204**, and an inverted version of the core clock signal. The BusEn2 signal **204** is a phase early since it is used for selecting a pulse for the bus clock prime signal **601**. To generate the bus clock, a latch **701** is used to latch the BusEn2 signal **204**, thereby delaying the BusEn2 signal **204** by a phase. Note that in this implementation, NAND gates are no longer used in favor of pass gates. Those skilled in the art would be familiar with the operation of such a circuit.

Information Transfer

To transfer information between clock domains (e.g., from the core clock domain to the bus clock domain, the present invention uses a series of enable signals and a state machine. These enable signals operate as synchronization signals for transfers. In one embodiment, these signals are referred to as write data enable, drive enable and sample enable. In one embodiment, the enable signals are generated using a state machine in the EBL that counts the number of core clocks until the next bus clock based on the bus clock ratio and reset.

The sample enable signal is used to transfer information from the bus clock domain to the core clock domain. FIG. 8A illustrates the use of the sample enable signal. Referring to FIG. 8A, information is transferred from the bus clock domain, through a pad cell **811**, and then to AND gate **801**. The AND gate **801** is also coupled to receive the sample enable signal **810**. When the sample enable signal **810** is high, the information at AND gate **801** is output to the core clock domain. In alternate embodiments, AND gate **801** may comprise a driver, an enable flip-flop, or any other device with the described functionality.

The write data enable and the drive enable signals are used to transfer information (data and signal information respectively) from the core clock domain. The write data enable is for write cycles when data is fetched from the data cache many coupled to the processor from a backside (dedicated) bus. The drive enable is a latched version of the write data enable used to transfer data from the core clock domain to the bus clock. FIG. 8B illustrates the use of the drive enable signal. Referring to FIG. 8B, information is transferred from the core clock domain to AND gate **802**. The AND gate **802** is also coupled to receive the drive enable signal **820**. When the drive enable signal **820** is high, the information at AND gate **802** is output through a pad cell **821** to the bus clock domain. In alternate embodiments, AND gate **802** may comprise a driver, an enable flip-flop, or any other device with the described functionality.

Similarly, FIG. 8C illustrates the use of the write data enable signal. Referring to FIG. 8C, information is transferred from the core clock domain to AND gate **803**. The AND gate **803** is also coupled to receive the write data enable signal **830**. When the write data enable signal **830** is

high, the information at AND gate **803** is output through a latch clocked by the core clock signal, which outputs to a pad cell **831** and then to the bus clock domain. In alternate embodiments, AND gate **803** may comprise a driver, an enable flip-flop, or any other device with the described functionality.

It should be noted that one skilled in the art would be able to generate such signals. Note that in FIGS. 8B and 8C, the AND gates are not required. All that is required is that information being transferred from the core clock domain is synchronized with the corresponding enable signal (e.g., each transferred signal is generated in the same clock as the enable signal).

FIG. 9 illustrates the clock and synchronization signals for the $\frac{2}{5}$ fractional speed bus. Referring to FIG. 9, the system clock signal **110**, the bus clock signal **202** and the core clock signal **201** are shown. Also shown in FIG. 9 is the bus clock prime enable signal **602** for the bus clock prime signal **601**. The bus clock prime signal **601** is generated by ANDing the core clock signal **201** with the bus clock enable signal **602**. While in the $\frac{2}{5}$ fractional bus case, the bus clock prime signal **601** alternates between $\frac{1}{2}$ and $\frac{1}{5}$ clock ratios, in the $\frac{2}{7}$ fractional bus case, it alternates between $\frac{1}{3}$ and $\frac{1}{4}$ clock ratios.

The enables signals are shown. In one embodiment, these signals are derived from the edge of the reset signal. The drive enable signal **820** is the delayed version of the write data enable signal **830**. In the $\frac{2}{5}$ fractional bus case, drive enable signal **820** is active at state **3**, one and a half clock cycles before the misaligned bus clock signal **202** (one clock before aligned). The sample enable signal **810** is active at state **1**, one half clock after the misaligned bus clock signal **202** (same clock during aligned). It should be noted that the sample enable signal **810** is not a latched version of the drive enable signal **820**. In fact, the sample enable signal **810** is either delayed from the drive enable signal **820** by one or two clock cycles on alternating cycles. Thus, in the $\frac{2}{5}$ fractional bus case, the sample enable signal **810** is asserted for a clock, deasserted for two clocks, asserted for a clock, deasserted for a clock, asserted for a clock, etc. A state machine uses a counter to count for every clock to identify the current state. The states **913** of the state machine are shown.

Note that the sample enable signal **810**, as well as the other enable signals, may be derived from one signal. In one embodiment, a series of flip-flops are coupled, with the output of each being the input to the next flip-flop. The write data enable signal **830** is input to the first flip-flop. The output of the first flip-flop is the drive enable signal **820**. The drive enable signal **820** is input into a second flip-flop. The output of the second flip-flop is one version of the sample enable signal **810** which is only one delay from the drive enable signal **820** (aligned with the bus clock signal). This version of the sample enable signal is also input into a third flip-flop, the output of which is the sample enable signal **810** with a two cycle delay from the drive enable signal **820** (misaligned with the bus clock signal). A multiplexor (mux) may then be used as a switch to select the correct sample enable signal based on whether the current bus clock signal is aligned or misaligned to the core clock. A signal from the state machine may be used to select the correct one of the sample enable signals as an output of the mux. In one embodiment, the determination of which to select first is based on whether the bus clock signal is aligned or misaligned with the core clock signal. An indication of which bus clock is misaligned is available is described below. Such a signal may be generated using a state machine synched at

11

reset. Thus, the sample enable signal **810** flips back and forth during aligned/misaligned periods.

An early system enable signal **912** and its enable **911** are also shown. In one embodiment, the early sample enable signal **912** is used by functional units to latch and generate a local sample enable. The enable **911** to create the early sample enable is used by the functional unit to generate the local sample enable signal **912**. The use of the early sample compensates for timing requirements (e.g., signal fanouts). Pads

FIGS. **10** and **11** illustrate the pad cell input path and output path respectively between the core clock domain and the pins. Referring to FIG. **10**, the pad cell input path **1000** comprises a series of latches. The pad cell comprises a pulsed latch **1001** coupled to receive input information from the pin. The output of this pulsed latch **1001** is coupled to the input of a latch **1002**. The output of latch **1002** is coupled to the core clock domain. In an alternative embodiment, the pulsed latch **1001** could be a D flip-flop.

The pulsed latch **1001** is clocked by an early version of the bus clock signal **202** (output signal from inverter **700** of FIG. **7**), while the latch **1102** is clocked by the core clock signal **201**. In another embodiment, the pulsed latch **1001** may be replaced by a flip-flop clocked by the bus clock signal **202**.

At the occurrence of a pulse of the bus clock signal **202**, the pulsed latch **1001** latches the information on the pin. If this rising edge of the bus clock signal **202** does not coincide with a rising edge of the core clock signal **201**, the data in the pulsed latch **1001** is sent to the core clock domain on the next high pulse of the core clock signal **201**. When the bus clock signal **202** and core clock signal **201** are misaligned (rising edges do not coincide), the latch **1002** acts to delay the transfer of information to the core for one half of a core clock cycle, in order to transfer this information from the bus clock time domain to the core clock time domain.

Referring to FIG. **11**, the cell output path of the present invention comprises a pad cell **1100** coupled between the core clock domain and the pin. The pad cell **1100** comprises a latch **1102** clocked by the core clock signal **201** that receives information from the core clock domain. The output of latch **1102** is coupled to the input of a master-slave flip-flop **1101** (e.g., D flip-flop), which is clocked by the bus clock signal **202**. The output of the flip-flop **1101** is coupled to the pin.

When information is being transferred from the core clock domain, the latch **1102** latches the information at the occurrence of a high phase of the core clock signal **201**. If this rising edge of the core clock signal **201** does not coincide with a rising edge of the bus clock signal **202**, the data in the latch **1102** is latched into flip-flop **1101** and is output to the pin at the next occurrence of the bus clock signal **202**. When the bus clock signal **202** and core clock signal **201** are misaligned (rising edges do not coincide), the latch **1102** acts to hold the data from the core for one half of a core clock cycle, in order to transfer this information from the core clock time domain to bus clock time domain.

In both the input and output paths, the extra latch clocked by the core clock signal **201** is necessary because of the misaligned core clock signal **201** and bus clock signal **202**. The data is held in the latch for one half of the core clock cycle in order to transfer it to the appropriate clock domain.

FIG. **12** illustrates one embodiment of pad cell implementation for the 2/N fractional bus. Referring to FIG. **12**, NAND gate **1201** is coupled to receive the signal going out from the core to the pad (which is connected to the pin), referred to in FIG. **12** as the coresigout signal **1250**, and a

12

drive-to-pin enable signal, referred to in FIG. **12** as the DRV2PIN signal **1251**. The output of NAND gate **1201** is coupled to the input of a master-slave flip-flop comprised of pass gates **1202** and **1205**, inverters **1203**, **1206**, **1207**, **1208**, **1209**, tristate buffer **1204** and NOR gate **1210**. The master in this master-slave flip-flop can be closed either on the rising edge of the core clock signal or the rising edge of the bus clock signal. This is accomplished through the use of NOR gate **1210**. The slave in this flip-flop is controlled by the bus clock signal. The operation of this flip-flop would be apparent to those skilled in the art. The output of the flip-flop is coupled to the input of inverter **1211**, the output of inverter **1211** is coupled to the input of output driver **1212** (implemented in one embodiment as a transistor), whose output is coupled to the pin. The pin is tied to a termination voltage through a termination resistor **1213**. The operation of this implementation would be well known to those skilled in the art.

The input path of the pad cell is shown in the lower half of FIG. **12**. An input signal is coupled to the sense amplifier **1216**. In an alternative embodiment, the sense amplifier **1216** may be a buffer. The output of the sense amplifier **1216** is coupled to the d input of the pulsed latch **1217** which is clocked by the early bus clock signal **110A**. The early bus clock signal **110A** is used to generate a pulse for the pulsed latch **1217** such that the timing matches that of a flip-flop clocked by the bus clock signal. In an alternative embodiment, the pulsed latch **1217** may be a flip-flop clocked by the bus dock signal.

The output of the pulsed latch **1217** is coupled to the input of a latch comprised of pass gate **1218**, inverters, **1219**, **1221**, **1222**, **1223** and tristate inverter **1220**. This latched is clocked by the core clock signal. The output of this latch is the signal into the core, referred to as sign2core signal **1252**.

In one embodiment, the core clock signal used to clock the core interface latches is replaced by the signals coreinclk **1221** and coreoutclk **1260**. There are specially designed signals to eliminate mindelay problems in the input and output paths.

With the addition of the latches to the core interface, a mindelay concern arises. The conditions occur when the core clock signal **201** and the bus clock signal **202** are out of phase (rising edges misaligned). These conditions are shown in FIG. **13**. For the path from core to pin, there is a concern that the latch driven on the coreoutclk signal will open too soon and change the data on the output (condition **1** in FIG. **13**). This transition from the coreoutclk signal must be either delayed or masked all together. For the path from the pin to the core, the concern is that the latch driven on the coreinclk signal will close too late and the data input to the core will change 1-half a core cycle too soon (condition **2**). To avoid this condition, the transition that closes the latch must occur sooner, or the opening of the latch for that phase be masked all together. Note that when the core clock signal **201** and bus clock signal **202** are aligned, the latches in the core clock domain need not be clocked and may be held open.

Referring to FIG. **13**, an inverted core clock signal **201** is shown with bus clock enable signals **203** and **204** that are used to generate the bus clock prime signal **602**. Also shown is an enable signal **1300** based on bus enable signal **204** to enable the out of phase of the bus clock signal **202**. Also demonstrated are the mindelay concerns outlined in the previous paragraph (conditions **1** and **2**). FIG. **13** shows the requirements for this signals to avoid mindelay problems. FIG. **15** illustrates the actual signals being used, including the coreinclk signal **1261** and coreoutclk signal **1260** used to

13

solve the mindelay problems in one embodiment of a 2/N pad cell. These signals only transition when needed; they do not toggle every core clock cycle.

FIG. 14 shows a circuit that may be used to generate the early bus clock signal **110a**, the bus clock signal **202**, the coreinclk signal **1261**, the coreoutclk signal **1260**, and the core clock signal **201**. The operation of this clock generation circuit would be well known to those skilled in the art, and will not be described in detail. The output waveforms generated by this circuit are shown in FIG. 15. Referring to FIG. 15, the rising edges of the coreinclk signal **1261** and coreoutclk signal **1260** match core clock signal timing. The timing of the falling edges for these signals is not critical. As shown in FIG. 15, the latch into the core is held open when the coreinclk signal **1261** is high, and the latch from the core to the pad is held open when the coreoutclk signal **1260** is low,

It should be noted that during 1/N modes (coreclk and busclk rising edges are always aligned), clocking of the 2/N interface latches are not necessary. In one embodiment, these latches are held open (not clocked) during 1/N modes. Logic Glitch

Due to the use of clocking signals of various domains, such as the core clock signal **201**, the bus clock signal **202** and the bus clock prime signal **601**, in conjunction of combinatorial logic, glitches sometimes occur. Specifically, the signals from three dock domains can qualify each other and cross over to the other domains if needed. For instance, the signals between the bus clock prime domain and the core clock domain may cross over easily because the core clock signal **201** and the bus clock prime signal **601** are always aligned. With respect to signals between the core clock domain and the bus clock domain, the pad cells described above operate to accommodate cases in which the core clock signal **201** and the bus clock signal **202** are aligned or misaligned. Signals transferred between the bus clock prime and the bus clock domains also do not have timing issues with respect to them because the core clock latch opening at the clock low state accommodates signals going from the bus clock prime domain to the bus clock domain and because the bus clock domain signals degenerate into the bus clock prime domain signal when going from the bus clock domain to the bus clock prime domain.

A glitch may occur due to the overlap between the bus clock signal **202** and the bus clock prime signal **601** in a misaligned clock. In other words, glitches occur when the bus clock prime information is qualified with bus clock information and the result is passed to the core clock domain without a synchronization signal (sample enable) and there is a misaligned core clock. To fix the glitch, an aligned clock signal, alignclksnn, is generated to mask out the glitch. In one embodiment, the alignclksnn signal is generated by a state machine. This signal is deasserted during misaligned clocks. Thus, the alignclksnn signal indicates when there are misaligned core clock edges between the bus clock signal **202** and the bus clock prime signal **601**.

FIG. 16 illustrates one embodiment of the use of the aligned clock signal. Referring to FIG. 16, a pad is coupled to a pad cell **1601** including the two latches **1601A** and **1601B** as described above. A signal (A) is output from latch **1601B** of the pad cell **1601** to one input of an AND gate **1603**. The other inputs of AND gate **1603** are coupled to the bus clock prime domain signal (B) from latch **1600** and the alignclksnn signal **1602**. The output (C) of AND gate **1603** is coupled to the input of latch **1604**, which is clocked by the core clock signal **201**. The output (D) of latch **1604** is a signal to the core clock domain. The alignclksnn signal **1602**

14

is used to qualify the result of ANDing signals A and B. A timing diagram illustrating a glitch and the occurrence of the alignclksnn signal **1602** is shown in FIG. 17 for the $\frac{2}{3}$ fractional speed bus. By deasserting the alignclksnn signal, during core clocks that are misaligned with the rising bus clock edge, as shown in FIG. 17, compensation for the glitch may be made.

In an alternate embodiments, the alignclksnn signal **1602** is not necessary in a situation such as that shown in FIG. 18A in which signal D is qualified by the drive enable signal **820** before it is propagated to the core clock domain. Referring to FIG. 18A, signals A and B output from the pad cell **1800** and the latch **1801** are coupled to inputs of AND gate **1802**. Latch **1801** is clocked by the bus clock prime signal **601**. The output (C) of AND gate **1802** is coupled to one input of AND gate **1806**, with the other input being coupled to the sample enable signal **810**. In this case, a glitch at state 2 will not be propagated to signal D (the output of AND gate **1806**) since the sample enable signal **810** is not asserted at state 2.

FIG. 18B illustrates another embodiment in which signal C, which is output from AND gate **1802**, is latch by latch **1803** which is clock by bus clock prime signal **601**. In this case, a glitch at state 2 will not be propagated at signal D (the output of latch **1803**) since the next rising edge of the bus dock prime signal **601** is at state 4 at which time the glitch is already gone.

Synchronization of Core Clock in 2/N Mode

Synchronization of the core clock signal may be required in the 2/N mode. Synchronization is required in FRC (Functional Redundancy Checking) mode where master and checker processors can power up with core clock signals out of phase. If the master and slave clocks are 180° out of phase in 2/N mode, an FRC error may result. That is, the two processors with opposite core clock phases will output the same data during different bus clock cycles, thereby resulting in FRC errors. Furthermore, synchronization may be required: in the testing environment. For instance, during silicon debugging, it is essential that the core clock signal and the bus clock signal be in a predetermined phase relationship to make tests reproducible every time. Particularly, it is desirable to have the core clock signal and bus clock signal to be in a predetermined phase at the time reset is deasserted. Although a reset event is used, any synchronization point may be used to identify when the bus and core clock signals are in phase. Because of these reasons, the core clock signal is synchronized to the bus clock in the 2/N mode. Note that synchronization is not needed in a 1/N mode since the bus clock signal and core clock signals are always in phase.

In the present invention, synchronization is performed with respect to the deassertion of the external reset signal. When reset is being deasserted, the phase relationship between the bus clock signal and the core clock signal is examined. If both are not in the high phase, the phase of the core clock signal is corrected to match that of the bus clock signal. Thus, the phase of the core clock signal is checked the time the reset is deasserted. If the core clock signal is in the low phase, then the core clock signal and bus clock signal are assumed to be in opposite phases. In this case, the core clock signal must be corrected to synchronize these clock signals. If the core clock signal is in the high phase, no action needs to be taken on the core clock signal to correct its clock phase. Note that in the present invention, correction is applied to the core clock signal only.

In an FRC mode, if the core clock signals of both processors are subjected to such correction, then the two processors will be synchronized once correction has com-

pleted. It is assumed that at the system bus clock signal is shared by both processors in this mode. In FRC mode, both processors must come up in the same state and at the same time after reset. In the 2/N clocking scheme of the present invention, when coming out of reset, there is no guarantee that the PLLs of both processors will be in sync such that the rising edges of their core clock signals are aligned with the system clock. The present invention provides for performing a set of events at a predetermined time after reset in order to synchronize the two PLLs. In the present invention, reset is held so that the processors have time to correct their clocks. Once this occurs, they will be properly synchronized out of reset.

Core Clock Correction

In one embodiment, the correction of the core clock signal is performed by shifting the core clock signal by one phase. Note that it is vital that the bus clock signal does not move or glitch because of the core clock correction. Otherwise, the PLL will go out of lock. Such a shifted pulse is shown in FIG. 19 where pulse 1981 is shown shifted to enable a correction. It should be noted that the shifting of the core clock signal (i.e., the correction) is performed after the system has released reset but prior to the time the processor leaves reset. Because the processor is still in reset, the remainder of the processor does not see the pulse extension to the core clock signal and, thus, does not do anything incorrect based on the extended clock pulse.

FIGS. 20 and 21 illustrate one embodiment of detector logic and associated waveforms, respectively, for a sync signal generator. The detector logic determines whether the phases of the bus and core clock signals are matching or not. Referring to FIG. 20, the detector logic 2000 comprises flip-flops 2001–2003, 2005–2007, inverters 2008 and 2009, latch 2004 and NOR gate 2010. Flip-flop 2001 is an inverting flip-flop. Flip-flops 2001–2003, latch 2004, and flip-flop 2005 are coupled in series, with the data output of one being coupled to the data input of the other, to generate the reset signal to the processor core. The data input of flip-flop 2001, the first in the series, is coupled to receive an active low reset signal from a pin. The output of flip-flop 2001 is the PaResetS00H signal and is fed to the inputs of flip-flops 2002 and 2006. Flip-flop 2005 generates the reset signal to the core of the processor. Flip-flops 2001–2003 are clocked by the bus clock signal, while latch 2004 and flip-flop 2005 are clocked by the core clock signal. Latch 2004 removes the phase path to the core because the signal is being transferred from a slow clock domain (the bus clock domain) to a fast clock domain (the core clock domain). Therefore, the use of latch 2006 eliminates race condition.

Flip-flops 2006 and 2007 along with inverter 2008 are coupled in series with the data input of flip-flop 2006 being coupled to the data output of flip-flop 2001. Flip-flop 2006 is clocked by the bus clock signal. The output of flip-flop 2006 (ResetS01H signal) clocks flip-flop 2007 which receives the core dock signal 201 on its data input. The output of flip-flop 2007 is coupled to the input of inverter 2008. The output of inverter 2008 is the core sync signal 2020 and is also coupled to the input of inverter 2009 which feeds back the core sync signal 2020 to one input of NOR gate 2010. The other input of NOR gate 2010 is the output of flip-flop 2006. The output of NOR gate 2010 sets flip-flop 2007.

The external reset signal is received by flip-flop 2001 in response to the bus clock signal 202 and is repeatedly “flopped” to set the correct timing for the reset signal to arrive at the processor core. During the time the reset signal is propagating through the serial chain of flip-flops and

latches 2001–2005, the reset is used as a clock to flip-flop 2007 after being provided by flip-flop 2006, which is docked by the bus dock signal 202. By being clocked in by the bus clock signal 202, the reset signal has the same, or coincident, phase with the bus clock signal 202. When the reset signal from flip-flop 2006 clocks flip-flop 2007, the core clock signal 201 is input into flip-flop 2007. If both the reset signal and the core clock signal 201 have the same phase, then the input “flopped” is the same as the reset signal; otherwise, it is different than the reset signal. If the bus clock signal 202 and the core clock signal 201 are mismatching in phase, then the core sync signal 2020 is zero. If they are matching in phase, the core sync signal 2020 is a one.

The feedback inverter 2009 in conjunction with NOR gate 2010 ensure that core sync signal 2020 preset low when the xreset# signal 900 at the pin is asserted (i.e., when the system goes into reset).

FIG. 21 illustrates the detector signal waveforms associated with the detector logic of FIG. 20. As shown in FIG. 21, the system reset is de-asserted (i.e., goes high), taking the system out of reset. When this occurs, the reset signal for the processor eventually goes high. However, prior to the processor exiting reset, the core clock signal 201 is sampled at the rising edge of the reset signal 2000 at flip-flop 2003 (the reset signal 502). When this occurs, the core sync signal 2020 goes high if there is a phase mismatch (solid line). If there is no phase mismatch, the core sync signal 2020 remains low (dotted line). Therefore, based on the state of the core sync signal 2020, the correction is made.

As discussed above, correction of the core clock signal 201 is performed by shifting the core clock signal 201 one phase. In one embodiment, in order to shift the core clock signal 201, the present invention uses the 2x frequency of the core clock signal 201 generated from the PLL. The PLL generates twice the frequency of the core clock signal 201 and is divide-by-two (via a divide-by-two counter) to obtain the 50% duty cycle core clock signal. Because the PLL is already generating a 2x frequency clock, the present invention generates a disable pulse to “swallow” one pulse of the 2x clock, thereby delaying one edge of the core clock signal 201 by a core clock phase in one cycle. FIG. 22A illustrates an example of the 2x clock having one of its pulses swallowed. By swallowing one cycle of the 2x frequency clock, a core clock phase is shifted 180°. By performing the shift, synchronization may be obtained.

FIG. 22B illustrates one embodiment of the circuitry and associated waveforms for shifting the core clock phase. The operation of the circuitry would be well-understood by those skilled in the art. The circuit outputs a zero until it detects core and bus clock signals mismatching in phase. At this point, the core sync signal 2020 goes from low to high, thereby causing a pulse of the 2x frequency clock from the PLL to be swallowed. In one embodiment, this circuitry detects not only the transition of the core sync signal from zero to one but also from one to zero to allow a return to an uncorrected condition. This may be necessary where the clock signals must be placed in a known state for testing.

In an alternate embodiment shown in FIG. 23, the core clock phase is shifted 180° at the rising edge of the reset signal if the core and bus clock signals are not in phase. As with the previous embodiment, the bus clock does not change, so the PLL remains locked. The 180° phase shift is achieved by selecting single inversion or double inversion of an early core clock. The operation of the circuitry in FIG. 23 would be well understood in the art. The early core clock signal is a core clock signal that is created prior to creation of the actual core clock signal to represent the core clock

signal when it is not present. This is described in greater detail below. Note that because the delay of the double inversion path cannot be designed to match the delay of the single inversion path in the clock tree over all process, temperatures and supply voltage corners, there will be some clock duty cycle variation between the two paths over the corners. It is conceivable, therefore, that some tests which pass when one path is selected will fail when the other is selected. This may result in test failures and is not reproducible every time. Because of the short coming, the previous implementation is more desirable.

Note that when the core clock is shifted, the present invention ensures that the bus clock signal **202** remains totally unchanged in phase and frequency; otherwise, the PLL will get out of lock and there will not be enough time for the PLL to relock before the integrated circuit starts executing commands. In the present invention, this is accomplished by clocking the counter that generates the bus clock enables by a separate core clock that does not change phase. If the core clock signal **201** has switched phase, the phase at which each of the two bus enables operates must also change so that they generate the bus clock from the correct core clock phase.

After performing a core clock correction, the bus clock enable signals are not accurate. That is, the BusEn1 signal that was selecting a high phase in the core clock signal is now selecting the core clock signal on the low phase. Likewise, the BusEn2 signal that was selected in a low phase of the core clock is now selecting in a high phase. In one embodiment, the bus enable signals can merely be swapped, particularly if it was only generating the bus clock signal. However, when generating the bus dock prime signal, swapping the two bus enable signals does not solve the problem of the bus clock enable signals. In one embodiment, to compensate for this, two sets of bus enables are generated, one that is valid if no correction is performed, while the other is valid if a correction is performed. Therefore, after correction has occurred, the present invention only changes the set of bus enables that are used.

One embodiment of the bus dock enable signal sets for before and after the shift is shown in FIG. 26. Referring to FIG. 26, the bus clock signal **2601** and the core dock signal **2602** are shown. The bus clock enable signal, BusEn1 and BusEn2, before the shift are shown as **2603** and **2604**, while those after the shift are **2605** and **2606**, respectively.

In an alternate embodiment, a shiftable divide-by-two generates a shifted clock. In the present invention, this shifted core clock signal is synchronized with the unshifted version such that when powering up, they power up the same way. In one embodiment, a shifted core clock is generated and a circuit reconstructs the unshifted clock from the shifted clock so that the two are always synchronized and their relationship is always correct. In one embodiment, a multiplexer may be used and selected based on the state of the core sync signal to either output the shifted clock or an inverted version of the unshifted clock. The unshifted version is used to supply a clock to the divide-by-N circuit described below. FIG. 24 illustrates such an arrangement. Testing of the Pulse Swallowing Circuitry

In order to test whether a pulse is actually swallowed, the present invention provides a test sequence for the pulse swallowing circuit. The use of the test sequence ensures that the logic that swallows a pulse is tested. More particularly, the present invention allows testing of cases of aligned and misaligned clock signals. The test sequence begins by identifying when the clocks will be aligned after asserting reset. Such an identification may be made by waiting a predeter-

mined number of cycles after reset has been deasserted so that if the clocks were misaligned, the core correction of the present invention would have produced an aligned clock signal. Any large number of cycles would suffice and would be well known to those skilled in the art. After the predetermined number of cycles, reset is asserted and deasserted again, and the device is either tested an even number of clocks or an odd number of clocks after the deassertion. After the testing, the reset is asserted and deasserted again, and the device is tested either an even number of clocks or an odd number of clocks, which ever number of clocks was not selected in the immediately preceding test. By testing in this way, there can be an assurance that the pulse swallowing circuitry of the present invention has been tested.

Reset Counter

In one embodiment, each processor is associated with an L2 cache memory that includes a PLL. During normal operation, the PLL of the L2 is in sync with the PLL of the processor. However, when using a non-integer ratio, there is no guarantee that both PLLs are in sync. Furthermore, by flipping the clock in the processor such as when a pulse is swallowed, synchronization between the PLL in the processor and the PLL in the L2 cache memory may be lost. The present invention provides for a technique to ensure that the two PLLs are in sync.

In the present invention, each PLL is typically comprised of an analog portion and a digital portion. The analog portion typically comprises the portion of the PLL that attempts to lock onto the incoming signal. The digital portion may be the component at the output of the PLL from which a feedback to the input of the PLL is produced. In one embodiment, this digital portion may comprise the pulse swallowing circuitry or a divide-by-N which produces an output and/or a feedback for the PLL. In the present invention, an output of a digital portion of the PLL of the processor provides the input to the PLL of the L2 cache memory. By doing so, any change in the output of the PLL of the processor may cause a similar change in the PLL of the L2 cache memory directly. To compensate for the additional time necessary for the L2 to regain lock, the present invention extends reset internally, thereby allowing the PLL of the L2 cache memory to sync up again. The reset is extended with a reset counter in the external bus control logic. Such a counter may be programmed and operates by preventing the processor from leaving reset until the counter counts out. In one embodiment, such a counter is always operating in this manner. However, such a counter may have to be disabled during testing.

Note that the reset counter of the present invention may be used to cause synchronization in FRC systems as well. One Embodiment of the 2/N Clocking Circuit

FIG. 25 is a block diagram of the 2/N clocking circuit of the present invention. Referring to FIG. 25, the PLL **2501** generates a 2x frequency signal as one input to pulse swallower **2502**. Pulse swallower **2502** also receives a core sync signal. Based on the core sync signal, pulse swallower **2502** either swallows a pulse of the clock signal to shift the core clock signal one phase or not. The output clock from pulse swallower **2502** is received by a divide-by-two **2503**.

The output of divide-by-two **2503** is an unshifted early clock signal **2504** and a shifted clock signal **2505**. The shifted early clock signal is coupled to dock distribution net **2508** and early bus and core clock generator **2509**. The unshifted early clock signal is coupled to divide-by-N **2510**, which outputs a signal to bus clock enable generation clock **2505**. The outputs of bus clock enable generation block **2511** are two bus clock enable signals which are sent to the clock

distribution network **2508** and the early bus and core clock generator **2509**.

Clock distribution network **2508** outputs the core clock signal to the processor and generates the bus clock signal from the core clock signal and the two bus clock enable signals. The bus clock signal is also provided to the PLL **2501**.

The early bus and core clock generator **2506** generates an early core clock signal and an early bus clock signal which are sent to the detector logic **2507**. In one embodiment, these are approximately 4 ns earlier than the regular bus and core clock signals. This early signal allows for shifting the core clock signal when the bus clock signal is not active. If the regular bus clock signal is used, there would not be enough time to perform the shifting. Detector logic **2507** generates the core sync signal to pulse swallower **2502** in response to the system reset signal (and a reset disable counter signal).

Whereas many alterations and modifications of the present invention will no doubt become apparent to a person of ordinary skill in the art after having read the foregoing description, it is to be understood that the various embodiments shown and described by way of illustration are in no way intended to be considered limiting. Therefore, references to details of various embodiments are not intended to limit the scope of the claims which in themselves recite only those features regarded as essential to the invention.

Thus, a method and apparatus for generating clock signals in an integrated circuit has been described.

What is claimed is:

1. A system comprising at least two agents wherein at least one of the agents comprises a clock generator that generates a core clock signal and a bus clock signal and wherein said at least one of the agents comprises:

a detector that samples the core clock signal upon the assertion of an externally-generated reset signal and determines whether the core clock signal is out of phase with the bus clock signal; and

correction circuitry coupled to the detector that places the core clock signal in phase with the bus clock signal, prior to the de-assertion of the reset signal, if the detector determines that the core clock signal is out of phase with the bus clock signal, such that the said at least two agents are synchronized after the de-assertion of the reset signal.

2. The system defined in claim 1 wherein the detector comprises:

a plurality of registers coupled to a reset pin;
a latch coupled in series to the plurality of registers;
a plurality of inverters coupled to the plurality of registers; and

a nor gate coupled to the plurality of inverters.

3. The system defined in claim 1 wherein a 2x frequency of the core clock signal is generated by a phase locked loop (PLL) is used to shift the core clock signal if the core clock signal is out of phase with the bus clock signal.

4. The system defined in claim 3 wherein a disable pulse is generated to swallow one pulse of the PLL.

5. The system defined in claim 1 wherein a single inversion of an early core clock signal is used to shift the core clock signal if the core clock signal is out of phase with the bus clock signal.

6. The system defined in claim 1 wherein a double inversion of an early core clock signal is used to shift the core clock signal if the core clock signal is out of phase with the bus clock signal.

7. The system defined in claim 1 wherein a shiftable divide-by-two circuit is used to shift the core clock signal if the core clock signal is out of phase with the bus clock signal.

8. A Functional Redundancy Checking (FRC) system comprising:

a clock source generating a system clock signal; and

a pair of processors responsive to the system clock signal; wherein each of the processors comprises:

a core clock generator generating a core clock signal;
bus enable generation logic to generate a first bus enable indication and a second bus indication; and
circuitry that generates a bus clock signal by selecting every N/2 cycles of the core clock signal in response to the first bus enable indication and the second bus enable indication, wherein N is an odd integer greater than two, such that the clock generator generates clock signals with a 2/N ratio of the bus clock signal to the core clock signal;

a detector that samples the core clock signal upon the assertion of an externally-generated reset signal and determines whether the core clock signal is out of phase with the bus clock signal; and

correction circuitry coupled to the detector that places the core clock signal in phase with the bus clock signal, prior to the de-assertion of the reset signal, if the detector determines that the core clock signal is out of phase with the bus clock signal, such that the pair of processors are synchronized after the de-assertion of the reset signal.

9. The method defined in claim 1 further comprising the steps of:

shifting the core clock signal to generate a shifted core clock signal;

unshifting the core clock signal to generate an unshifted core clock signal for the shifted core clock signal; and

selecting either the shifted core clock signal or the unshifted core clock signal based on whether the core clock signal is out of phase or in phase with the bus clock signal respectively.

10. The method defined in claim 1 wherein the step of checking the phase of the core clock signal comprises the steps of:

sampling the core clock signal at the rising edge of a reset signal;

detecting whether the core clock signal is out of phase with the bus clock signal; and

asserting a mismatch indication if the core clock signal is detected as being out of phase.

11. The method defined in claim 1 wherein the core clock signal is generated by phase-locked loop (PLL) outputting a frequency clock to a divide-by-two counter, and further wherein the step of correcting the core clock signal comprises the step of preventing one cycle of the frequency clock from reaching the divide-by-two counter.

12. The method defined in claim 1 wherein the bus clock signal is generated using a plurality of bus clock enable indications, and further comprising the step of switching two of the bus clock enable indications having pulses on opposite core clock phases, such that cycles of the core clock signal selected by the two bus clock enable indications are opposite in phase to cycles selected prior to switching.

13. The method defined in claim 1 further comprising the steps of:

generating two sets of bus clock enable indications for use in generating the bus clock signal; and

selecting one of the two sets of bus clock enable indications based on whether the core clock signal undergoes correction.

21

14. The method defined in claim 1 wherein the core clock signal is shifted by inverting a predecessor version of the core clock signal to attain a 180° phase shift in the core clock signal.

15. A method for internally synchronizing to a 2/N mode 5 core clock signal with a bus clock signal within an integrated circuit, said method comprising:

22

checking the phase of the core clock signal when a reset signal is de-asserted; and
correcting the core clock signal by shifting the core clock signal by one phase to match the bus clock if the core clock signal is out of phase with the bus clock signal.

* * * * *

UNITED STATES PATENT AND TRADEMARK OFFICE
CERTIFICATE OF CORRECTION

PATENT NO. : 6,208,180 B1
DATED : March 27, 2001
INVENTOR(S) : Fisch et al.

Page 1 of 1

It is certified that error appears in the above-identified patent and that said Letters Patent is hereby corrected as shown below:

Column 20,

Lines 27, 38, 47, 53 and 60, delete “in claim **1**” and insert -- in claim **8** --.


Column 21,

Line 1, delete “in claim **1**” and insert -- in claim **8** --.

Signed and Sealed this

Fifteenth Day of October, 2002

Attest:

A handwritten signature in black ink, appearing to read "James E. Rogan", with a long horizontal flourish extending from the bottom of the signature.

Attesting Officer

JAMES E. ROGAN
Director of the United States Patent and Trademark Office